



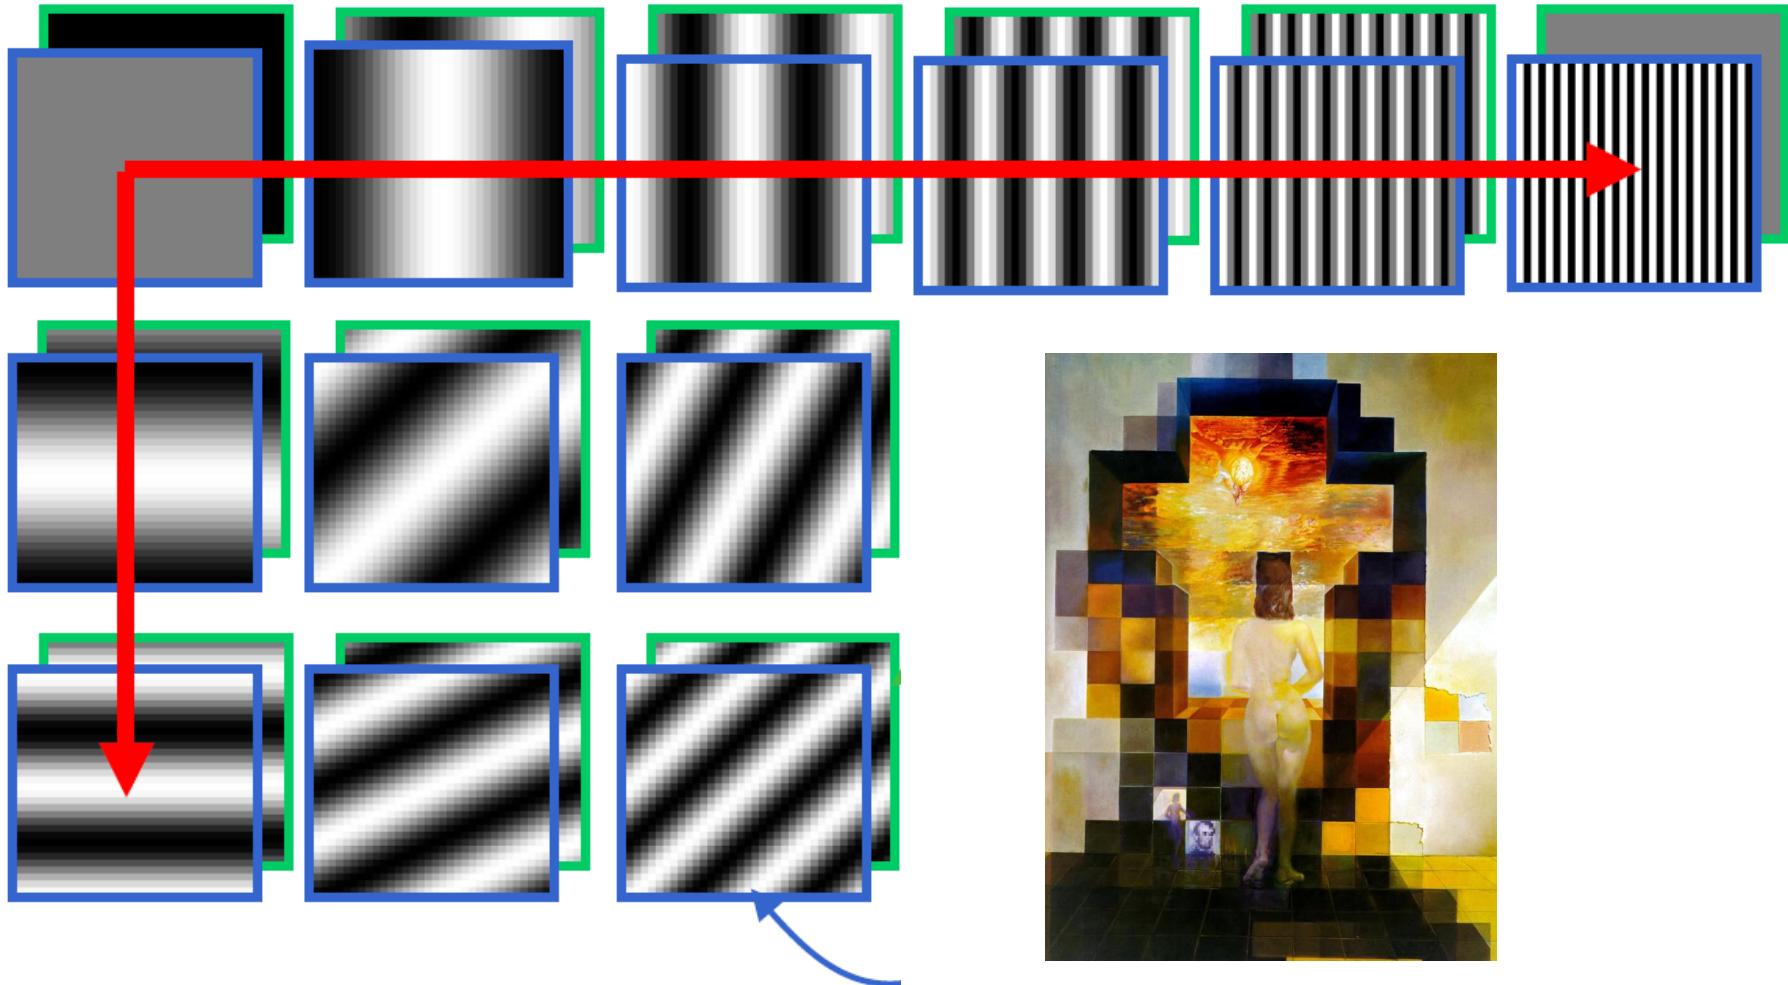
# Recap: Fourier domain

# 2d Fourier Transform

$$\hat{F}(k, \ell) = \frac{1}{\sqrt{MN}} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} F(m, n) e^{-j2\pi \left( k \frac{m}{M} + \ell \frac{n}{N} \right)}$$

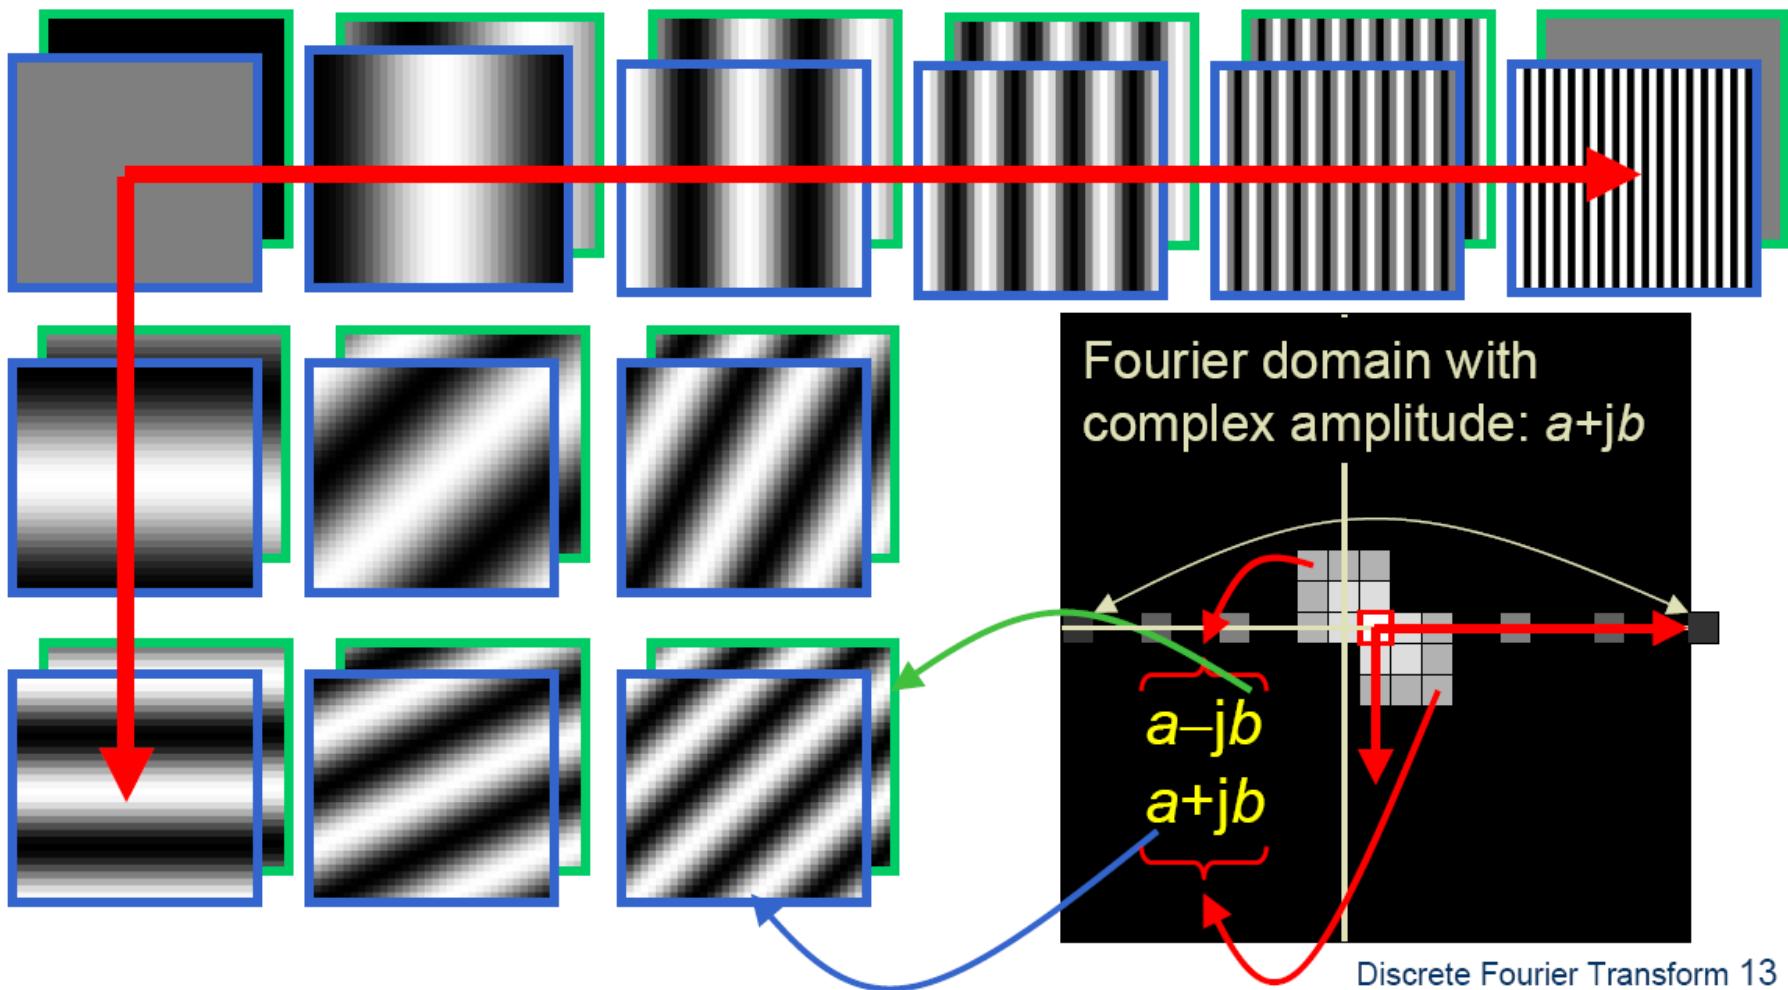
# Fourier Bases

Teases away fast vs. slow changes in the image.



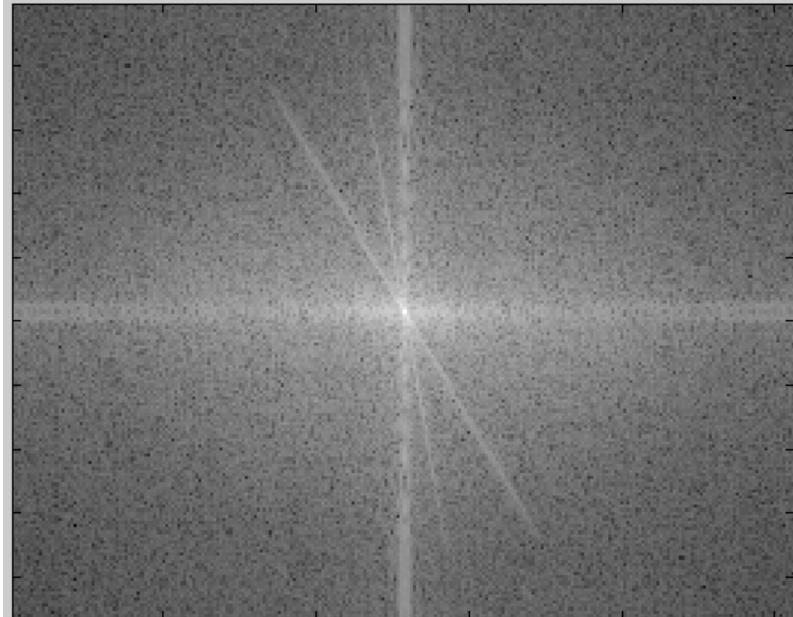
This change of basis is the Fourier Transform

# Fourier Bases

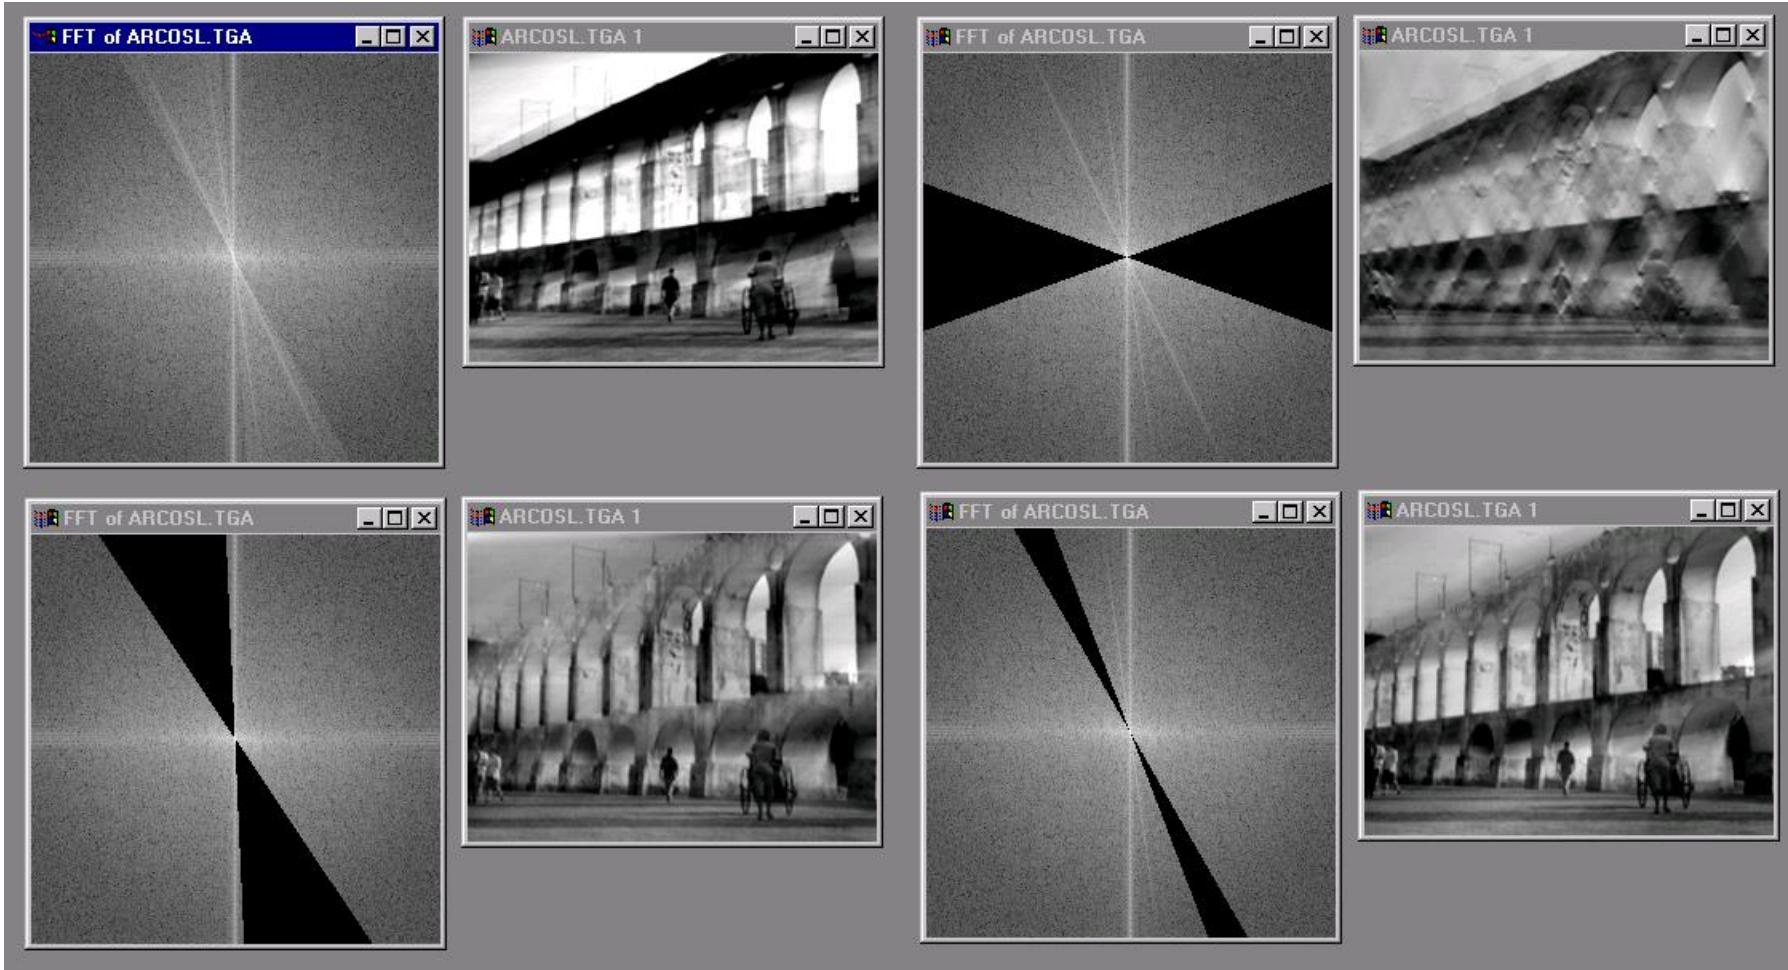


in Matlab, check out: `imagesc(log(abs(fftshift(fft2(im)))));`

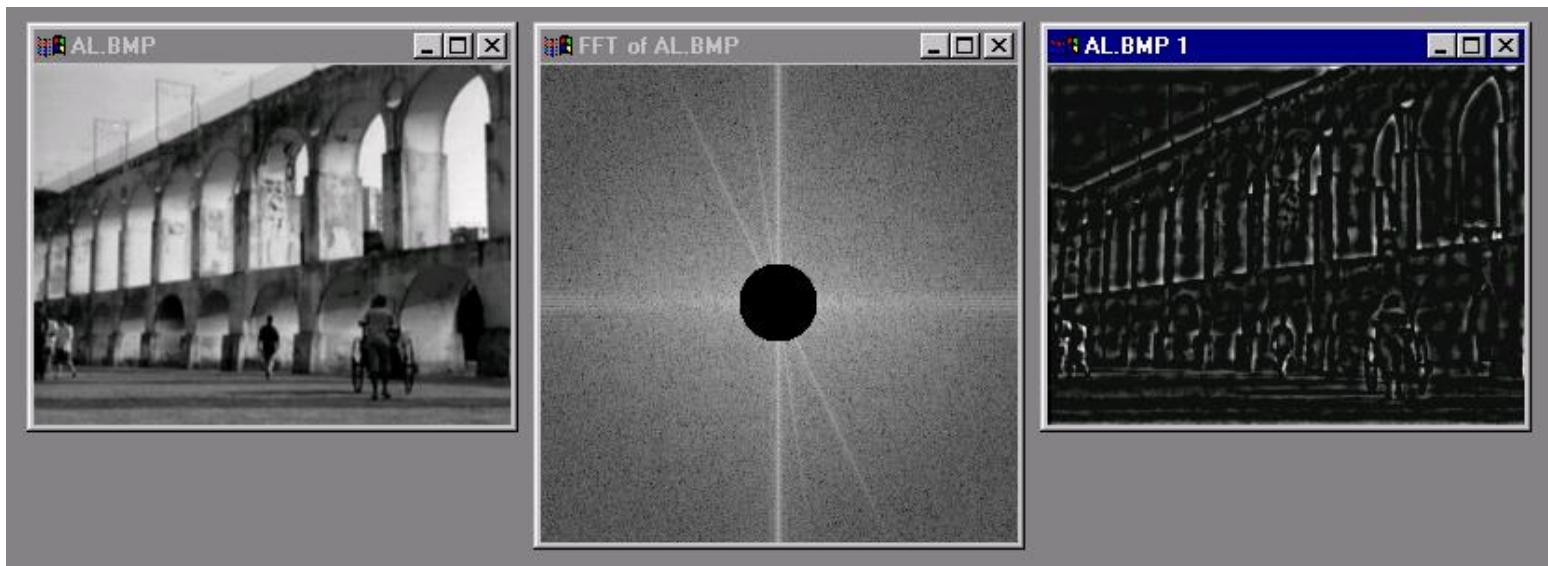
# Man-made Scene



# Can change spectrum, then reconstruct

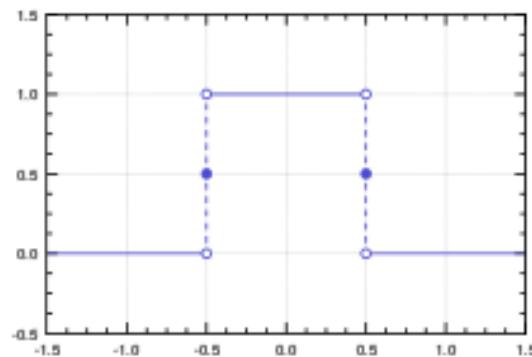


# Low and High Pass filtering

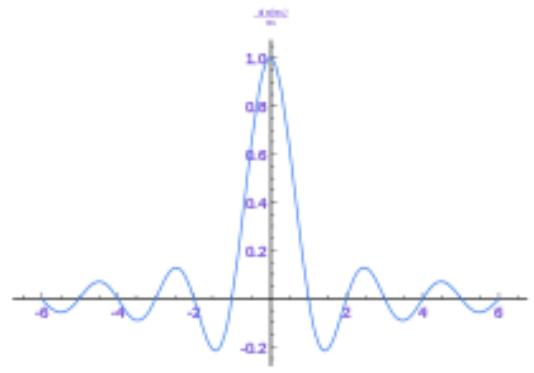


# Sinc Filter

- What is the spatial representation of the hard cutoff in the frequency domain?



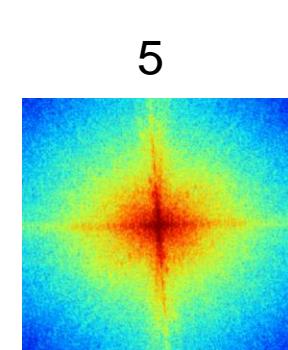
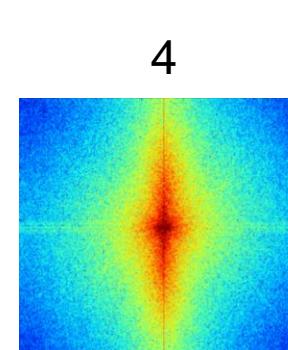
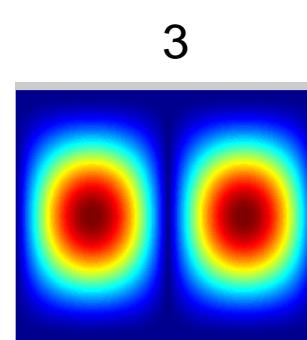
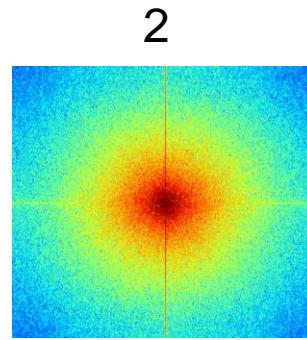
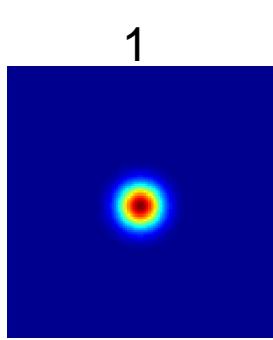
Frequency Domain



Spatial Domain

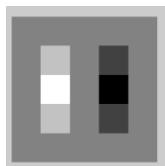
# Review

1. Match the spatial domain image to the Fourier magnitude image



B

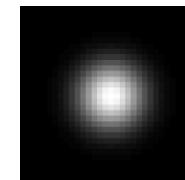
A



C



D



E



# The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

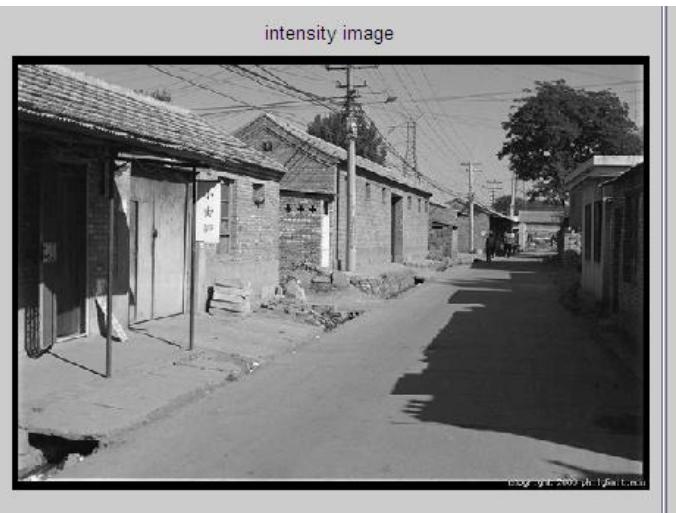
$$F[g * h] = F[g]F[h]$$

- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

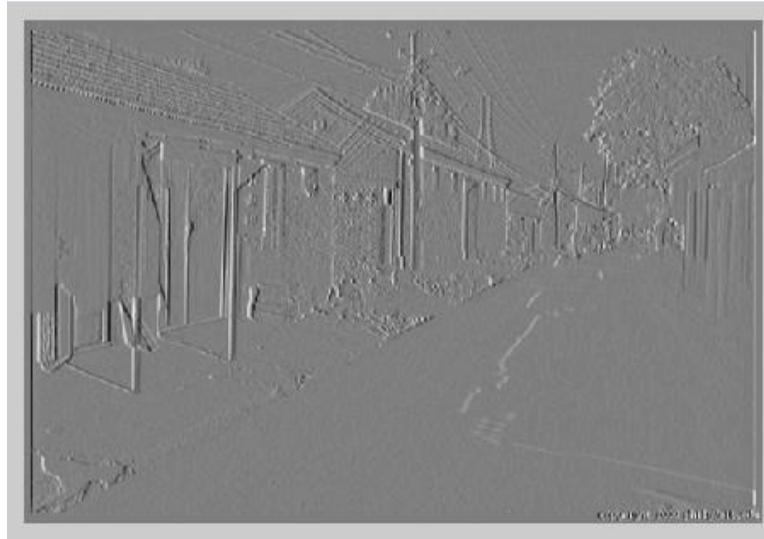
$$g * h = F^{-1}[F[g]F[h]]$$

# Filtering in spatial domain

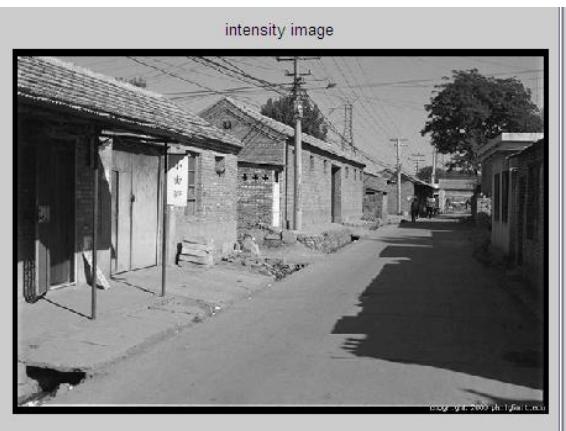
1	0	-1
2	0	-2
1	0	-1



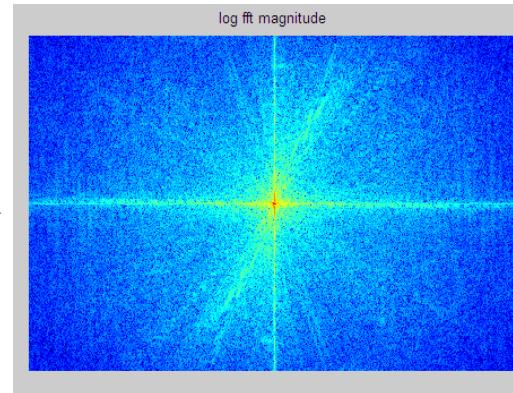
$$\ast =$$



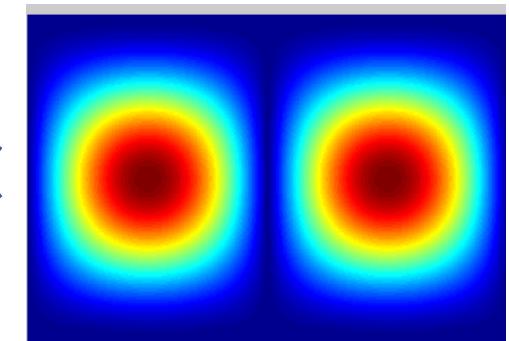
# Filtering in frequency domain



FFT  
→

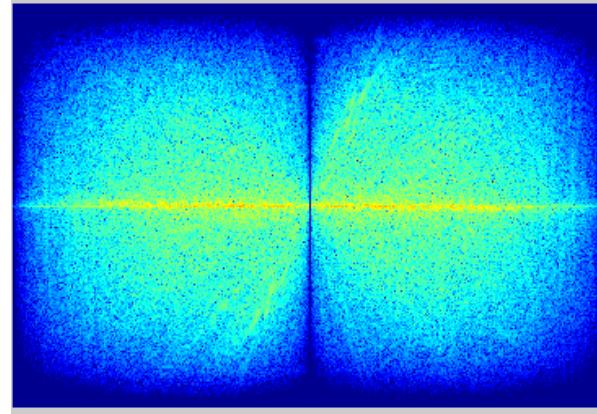


×



||

Inverse FFT  
←



# Fourier Matlab demo

# FFT in Matlab

- Filtering with fft

```
im = double(imread('...'))/255;
im = rgb2gray(im); % "im" should be a gray-scale floating point image
[imh, imw] = size(im);

hs = 50; % filter half-size
fil = fspecial('gaussian', hs*2+1, 10);

fftsize = 1024; % should be order of 2 (for speed) and include padding
im_fft = fft2(im, fftsize, fftsize); % 1) fft im with padding
fil_fft = fft2(fil, fftsize, fftsize); % 2) fft fil, pad to same size as
image
im_fil_fft = im_fft .* fil_fft; % 3) multiply fft images
im_fil = ifft2(im_fil_fft); % 4) inverse fft2
im_fil = im_fil(1+hs:size(im,1)+hs, 1+hs:size(im, 2)+hs); % 5) remove padding
```

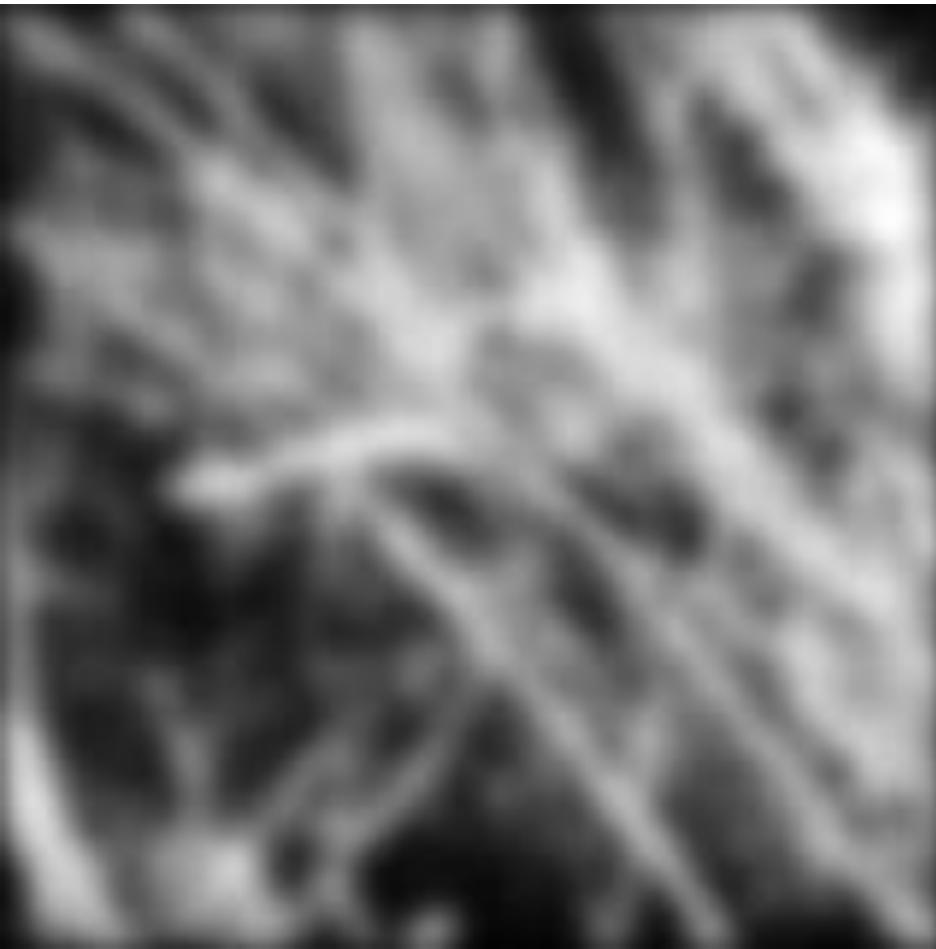
- Displaying with fft

```
figure(1), imagesc(log(abs(fftshift(im_fft))), axis image, colormap jet
```

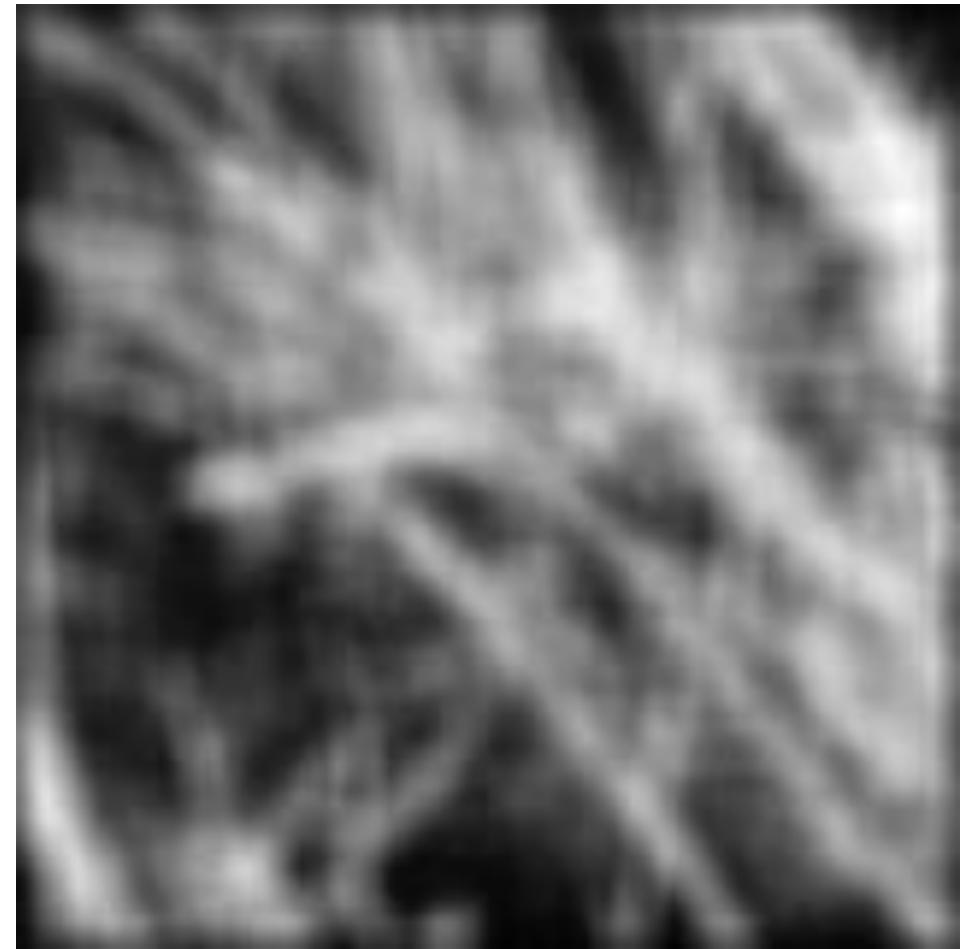
# Filtering

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

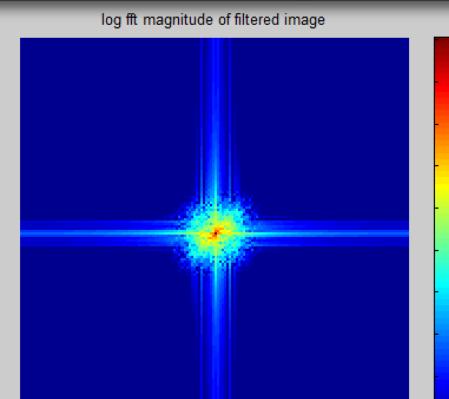
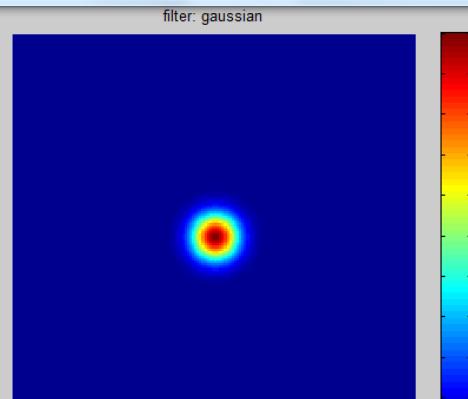
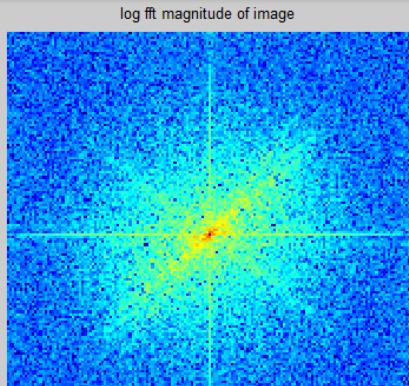
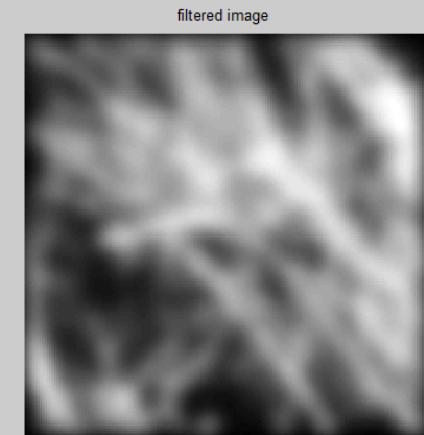
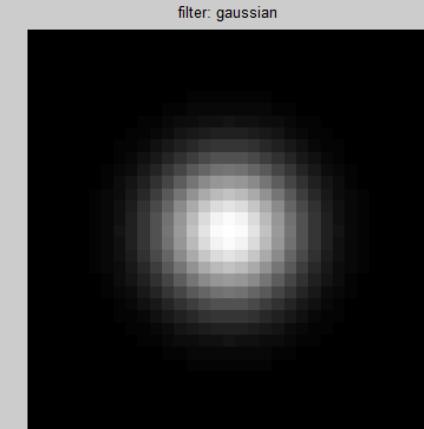
Gaussian



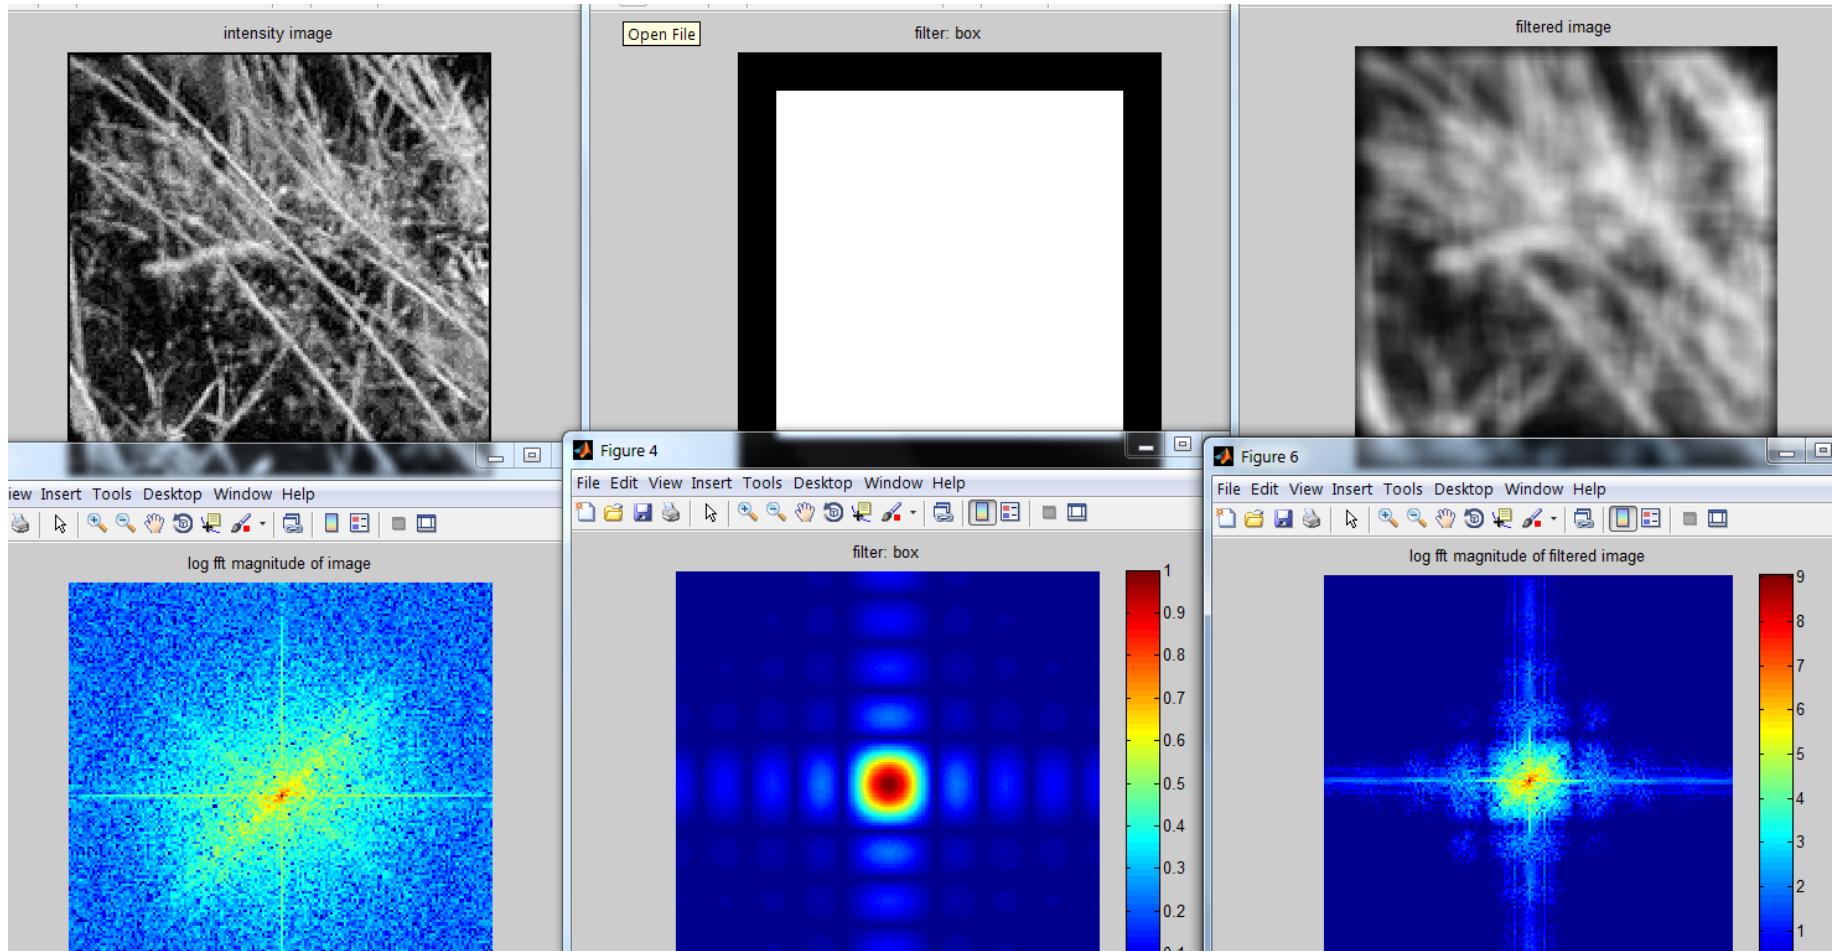
Box filter



# Gaussian



# Box Filter



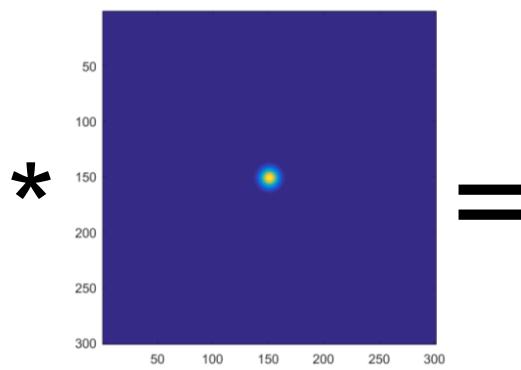
# Is convolution invertible?

- If convolution is just multiplication in the Fourier domain, isn't deconvolution just division?
- Sometimes, it clearly is invertible (e.g. a convolution with an identity filter)
- In one case, it clearly isn't invertible (e.g. convolution with an all zero filter)
- What about for common filters like a Gaussian?

# Let's experiment on Novak



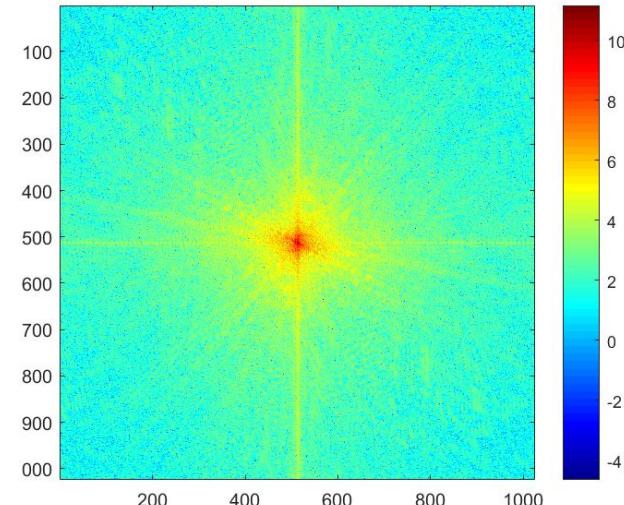
# Convolution



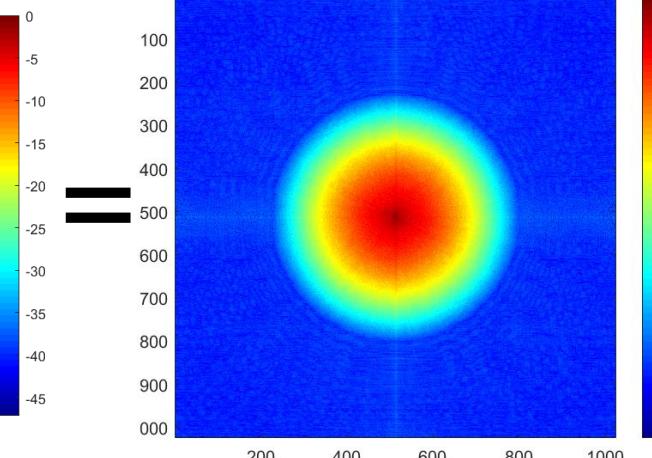
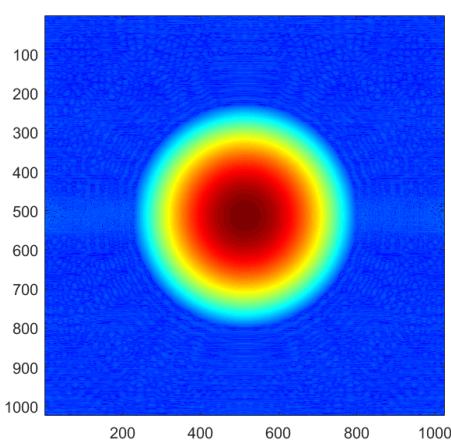
FFT

FFT

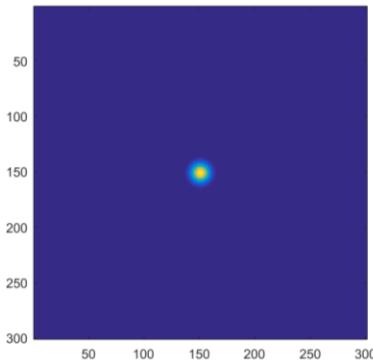
iFFT



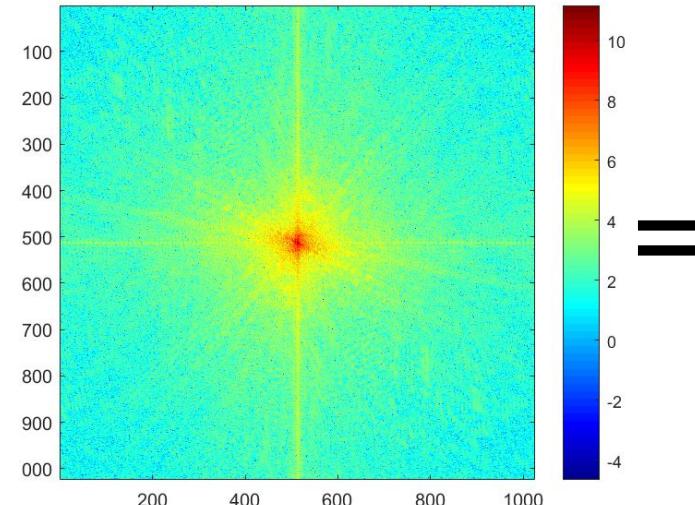
\*



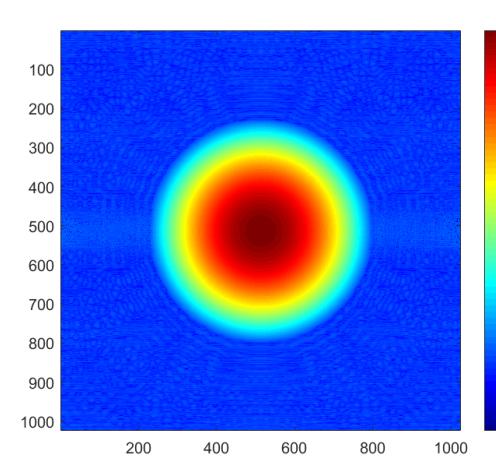
# Deconvolution?



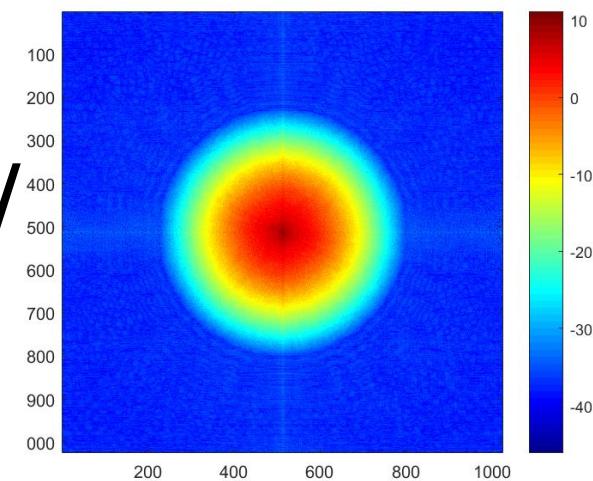
iFFT



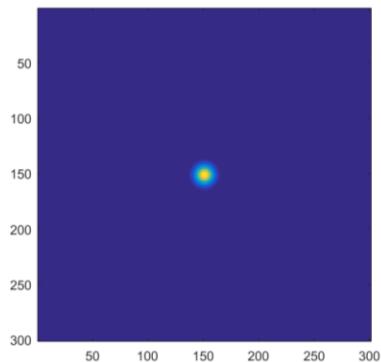
FFT



FFT



# But under more realistic conditions



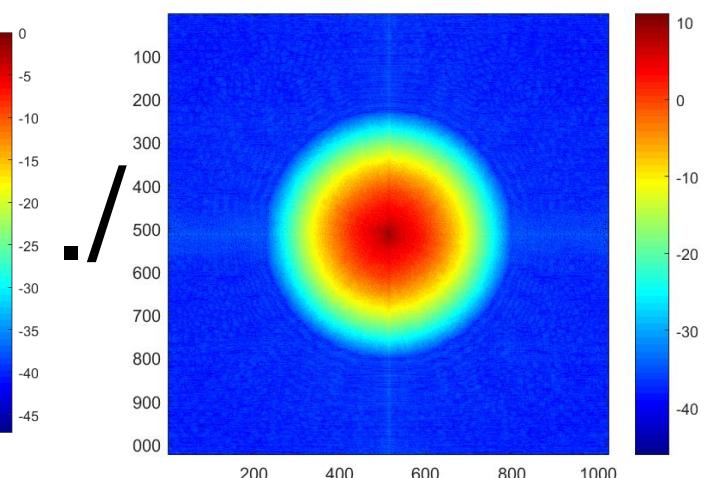
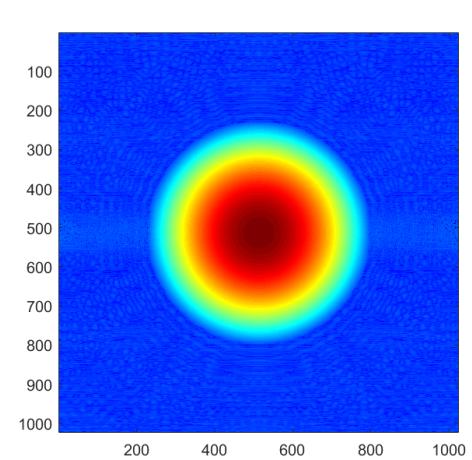
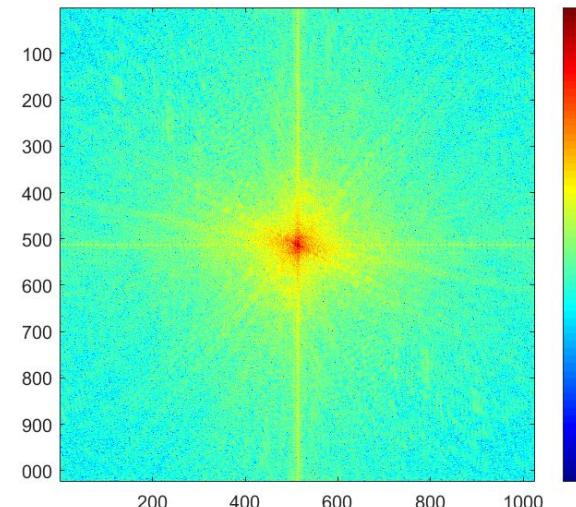
Random noise, .000001 magnitude



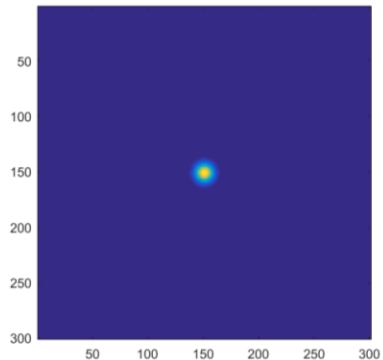
iFFT

FFT

FFT



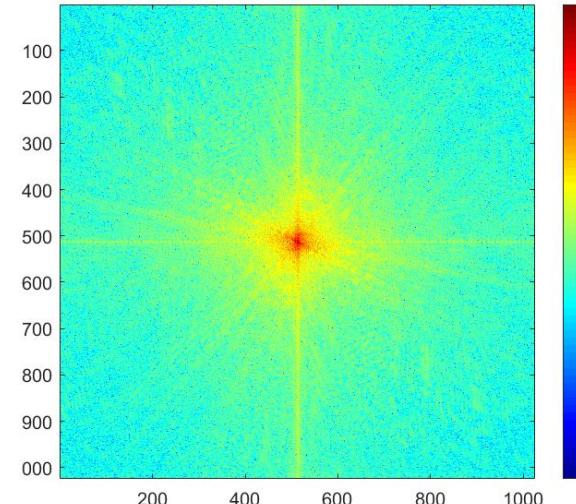
# But under more realistic conditions



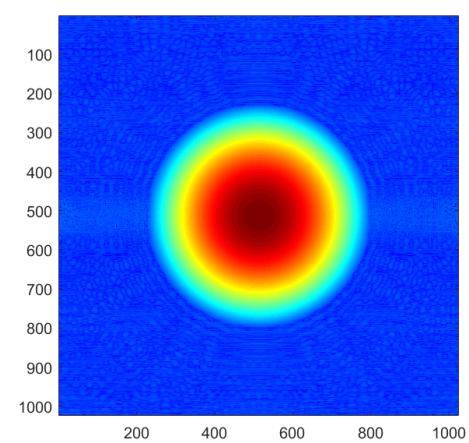
Random noise, .0001 magnitude



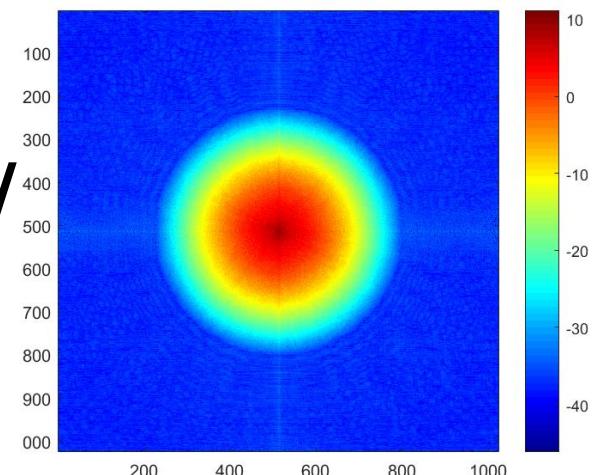
iFFT



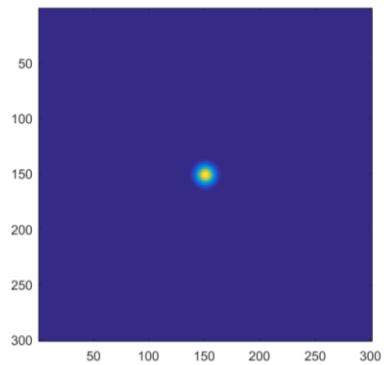
FFT



FFT



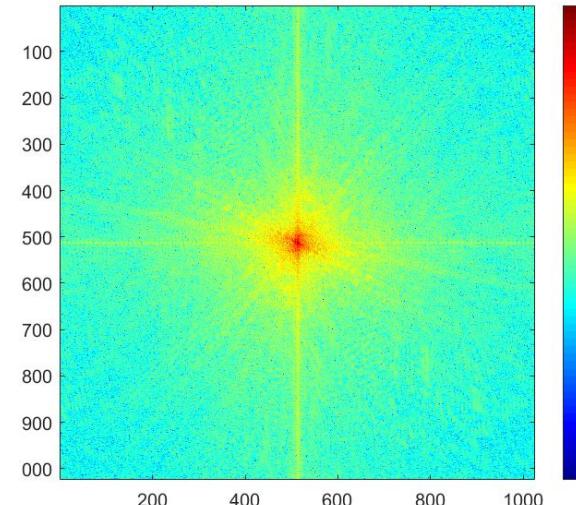
# But under more realistic conditions



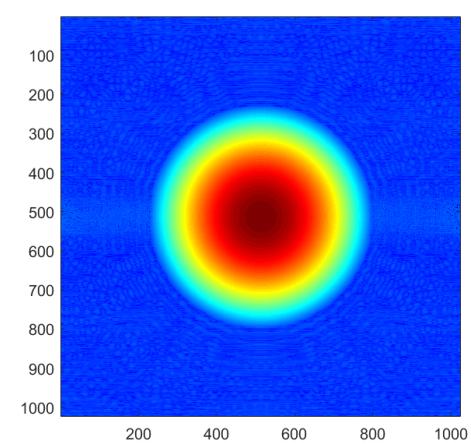
Random noise, .001 magnitude



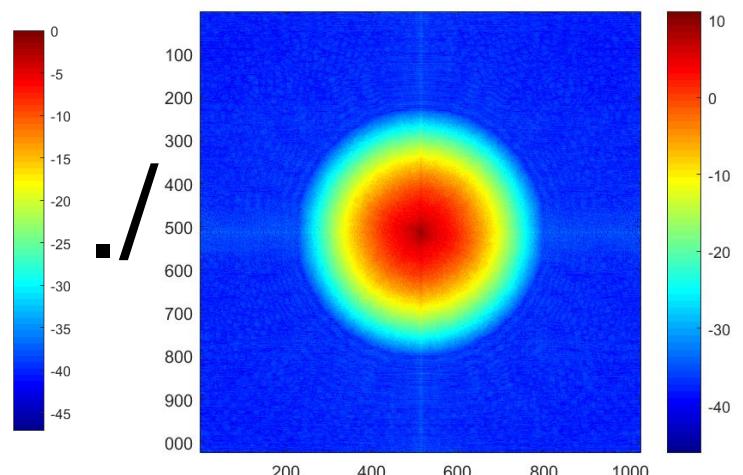
iFFT



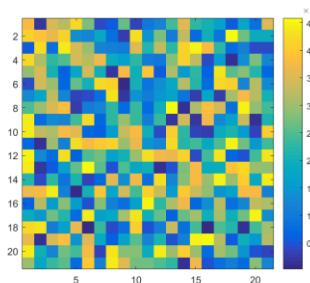
FFT



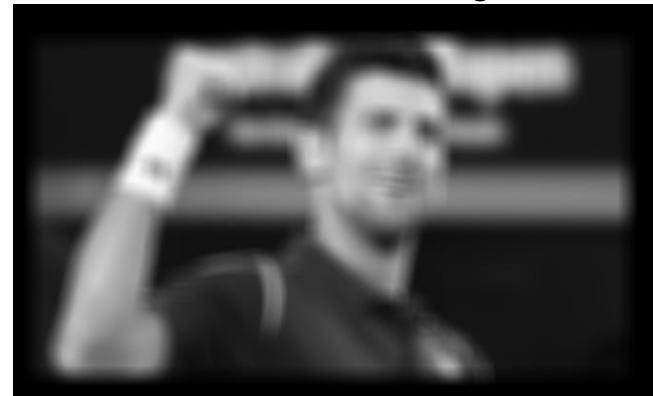
FFT



# With a random filter...



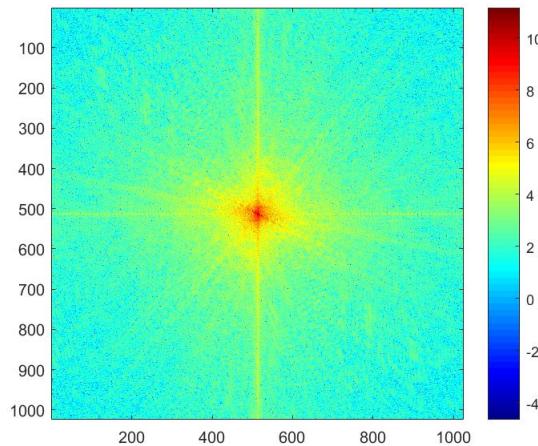
Random noise, .001 magnitude



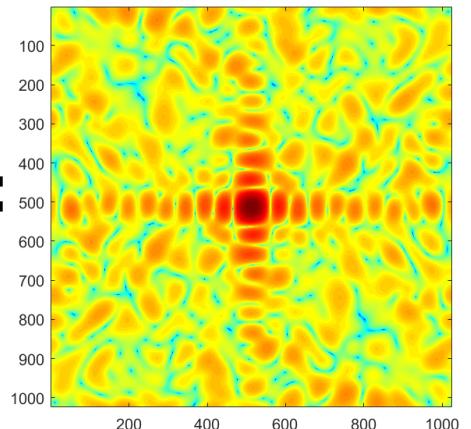
iFFT

FFT

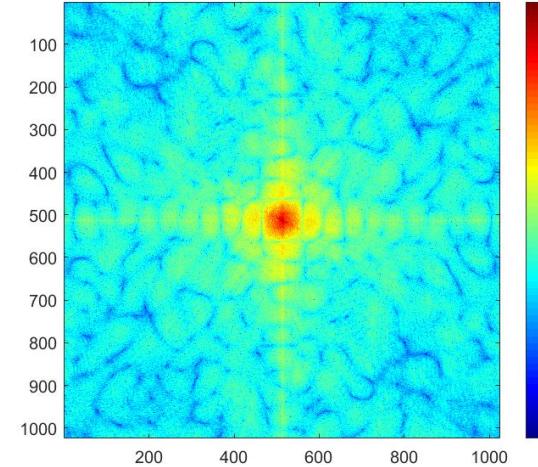
FFT



=



/

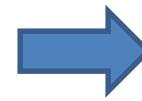


# Deconvolution is hard

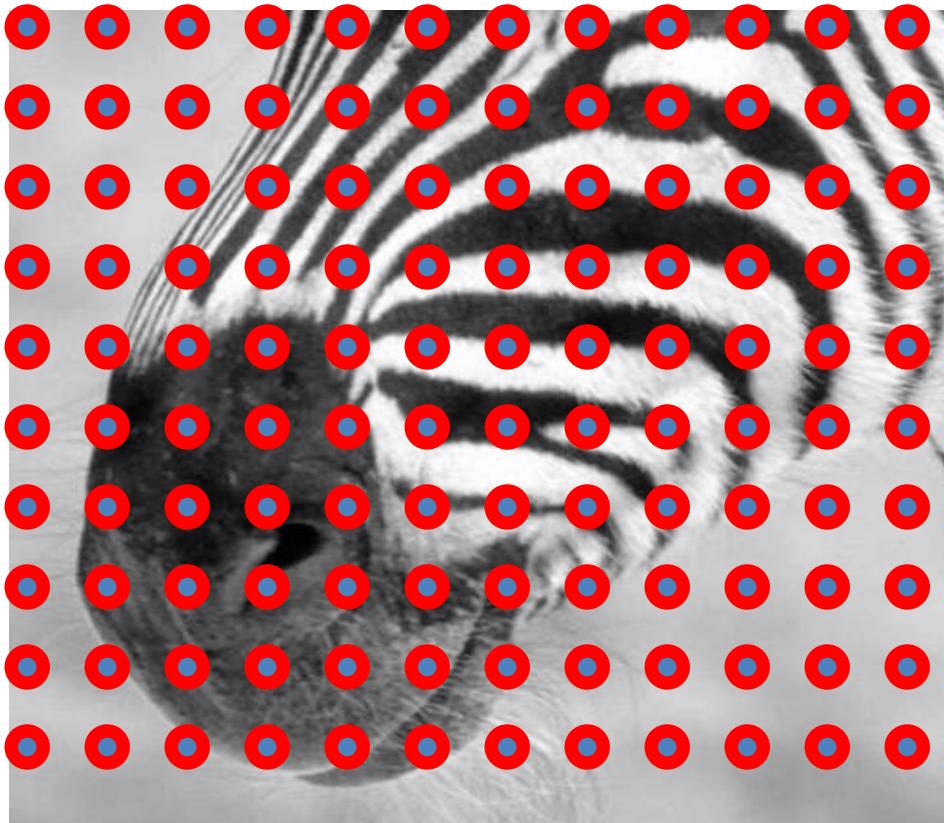
- Active research area.
- Even if you know the filter (non-blind deconvolution), it is still very hard and requires strong *regularization*.
- If you don't know the filter (blind deconvolution) it is harder still.

# Sampling

**Why does a lower resolution image still make sense to us? What do we lose?**



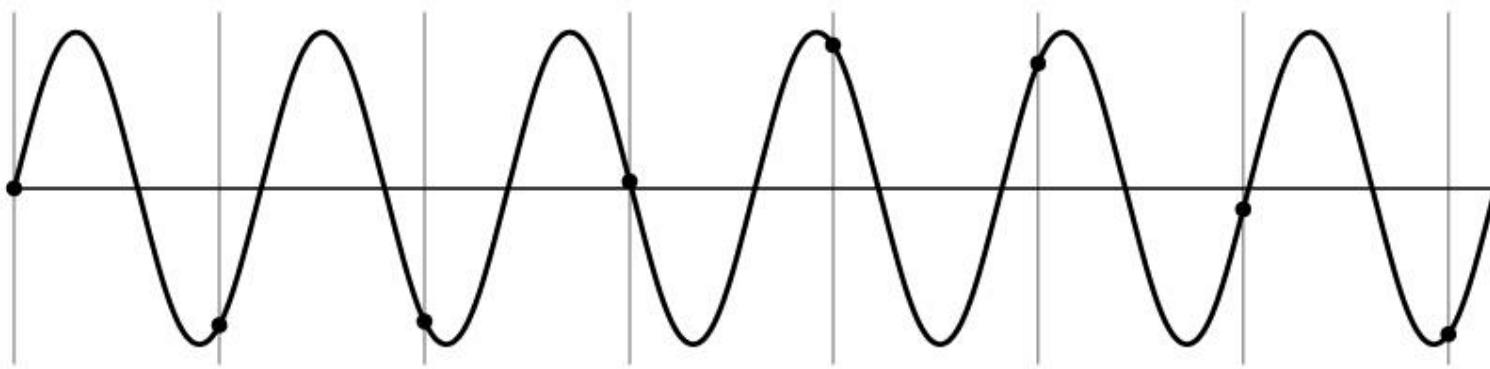
# Subsampling by a factor of 2



Throw away every other row and column to create a 1/2 size image

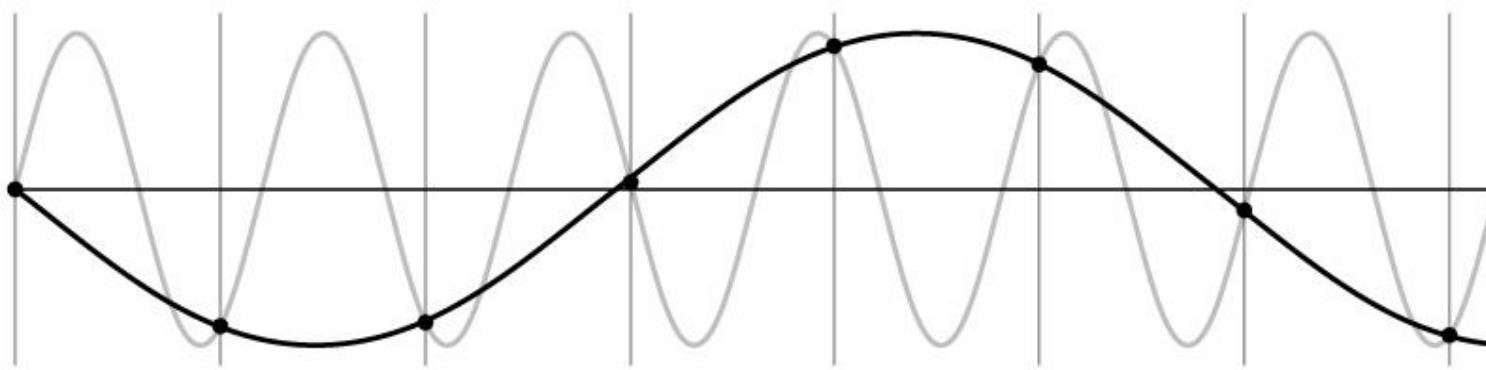
# Aliasing problem

- 1D example (sinewave):



# Aliasing problem

- 1D example (sinewave):



# Aliasing problem

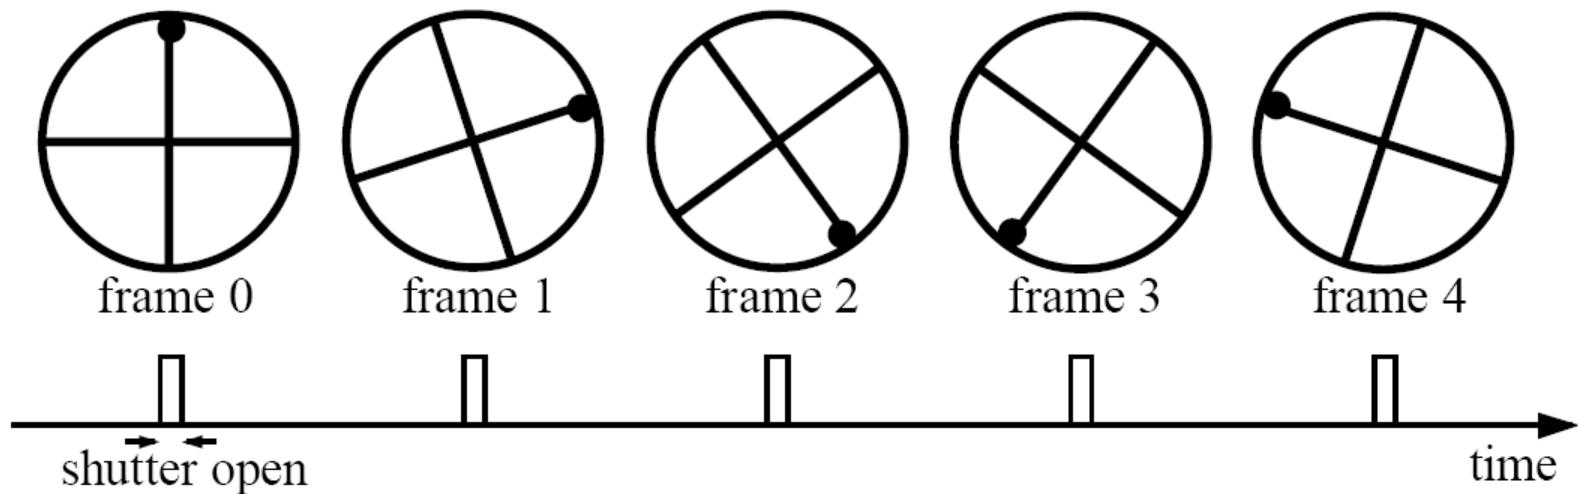
- Sub-sampling may be dangerous....
- Characteristic errors may appear:
  - “car wheels rolling the wrong way in movies”
  - “Checkerboards disintegrate in ray tracing”
  - “Striped shirts look funny on color television”

# Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time =  $1/30$  sec. for video,  $1/24$  sec. for film):



Without dot, wheel appears to be rotating slowly backwards!  
(counterclockwise)

# Aliasing in graphics



# Sampling and aliasing

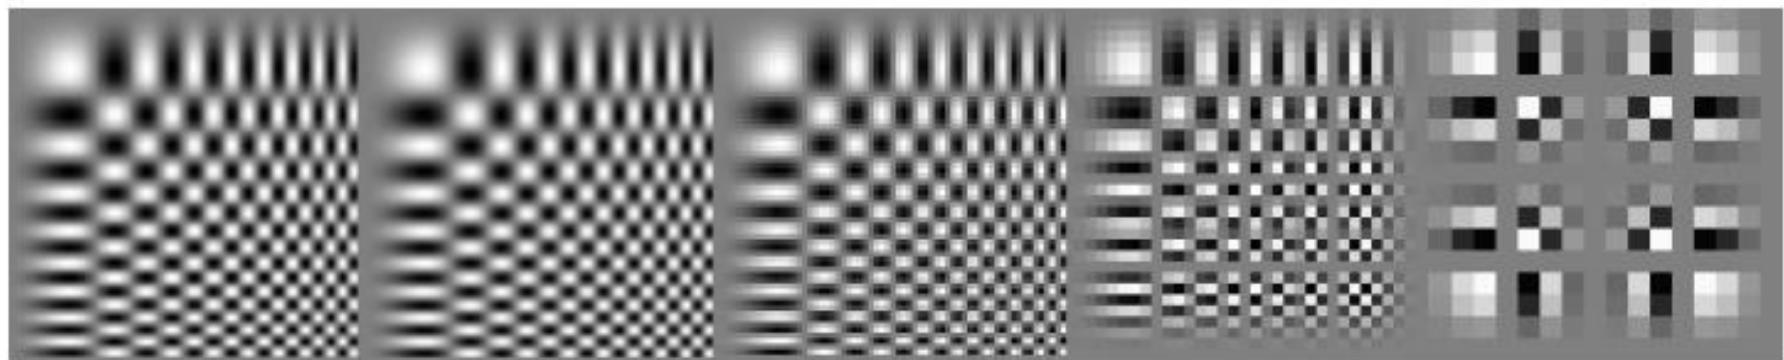
256x256

128x128

64x64

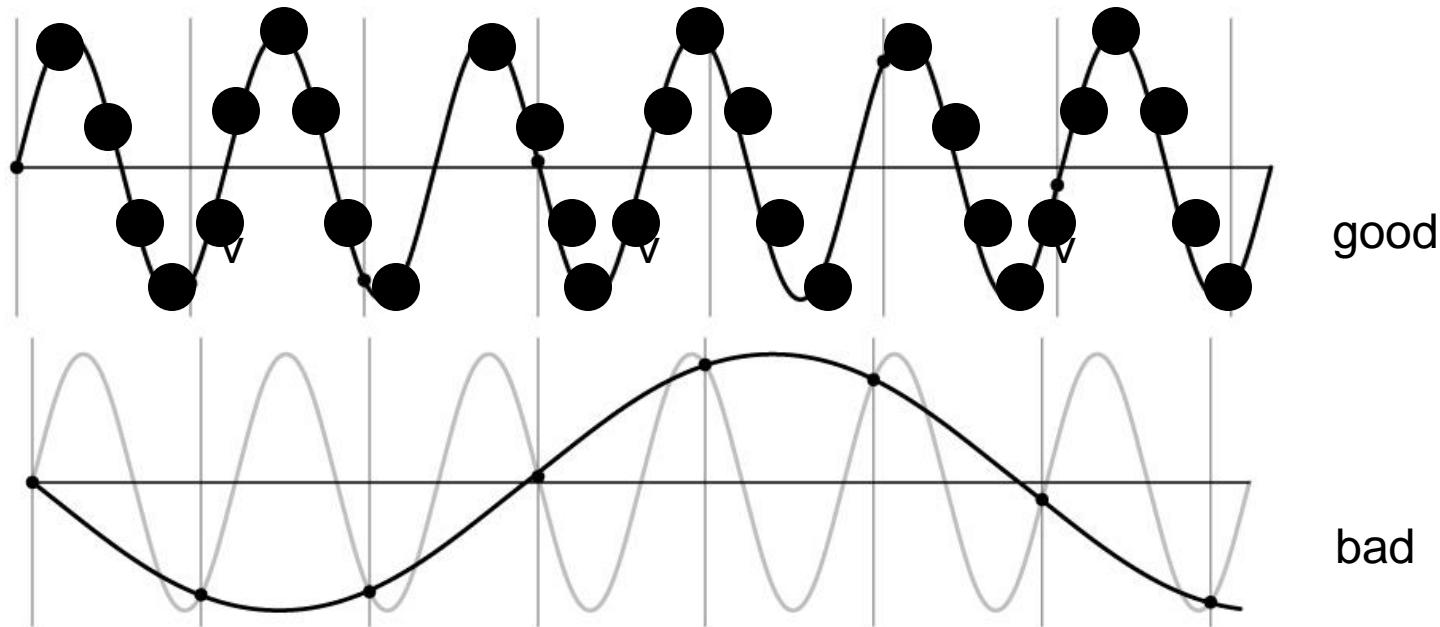
32x32

16x16



# Nyquist-Shannon Sampling Theorem

- When sampling a signal at discrete intervals, the sampling frequency must be  $\geq 2 \times f_{\max}$
- $f_{\max}$  = max frequency of the input signal
- This will allow to reconstruct the original perfectly from the sampled version



# Anti-aliasing

Solutions:

- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
  - Will lose information
  - But it's better than aliasing
  - Apply a smoothing filter

# Algorithm for downsampling by factor of 2

1. Start with image( $h, w$ )

2. Apply low-pass filter

```
im.blur = imfilter(image, fspecial('gaussian', 7, 1))
```

3. Sample every other pixel

```
im.small = im.blur(1:2:end, 1:2:end);
```

# Anti-aliasing

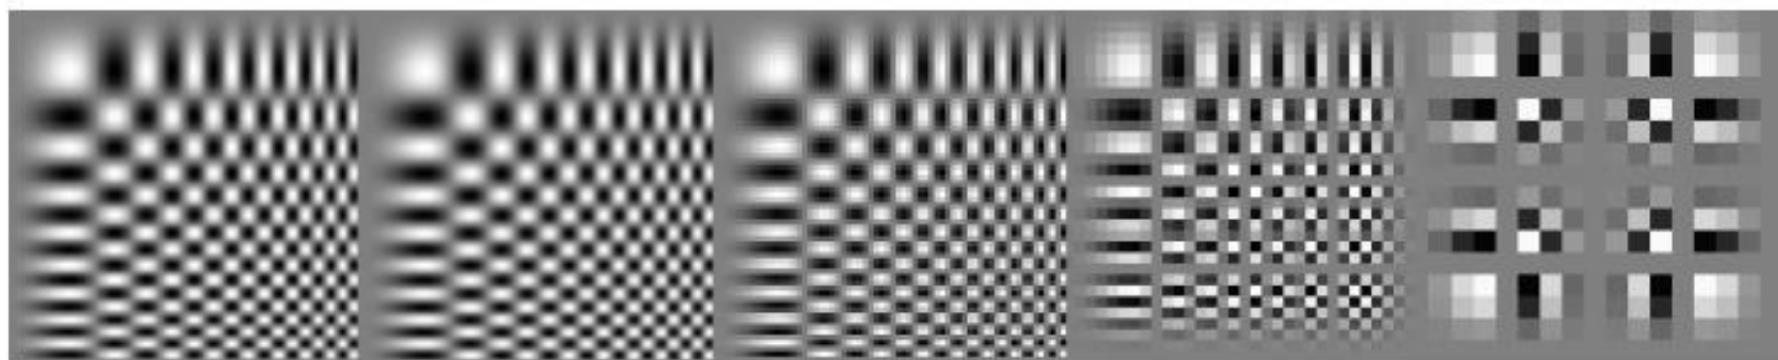
256x256

128x128

64x64

32x32

16x16



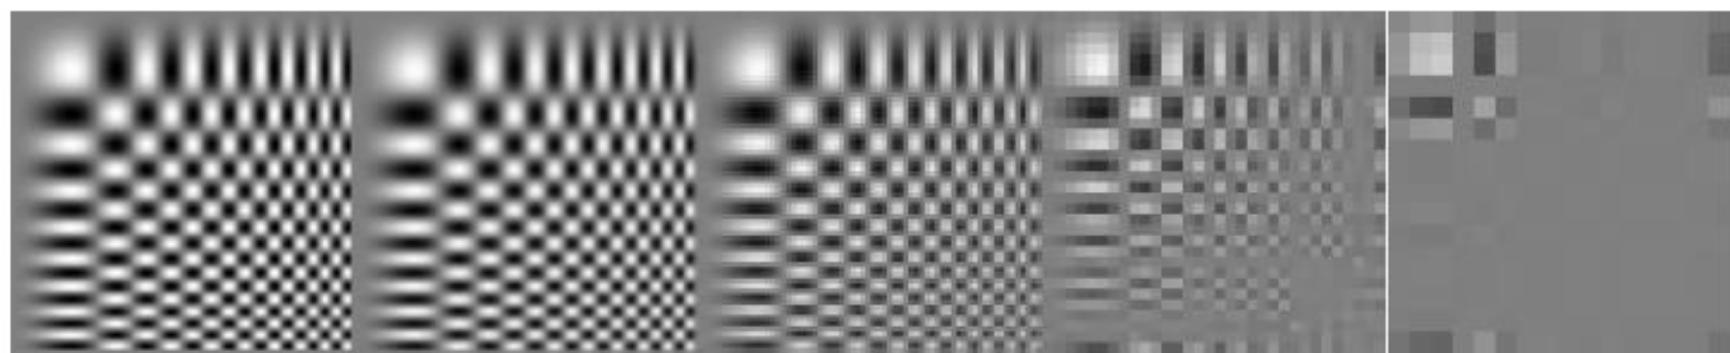
256x256

128x128

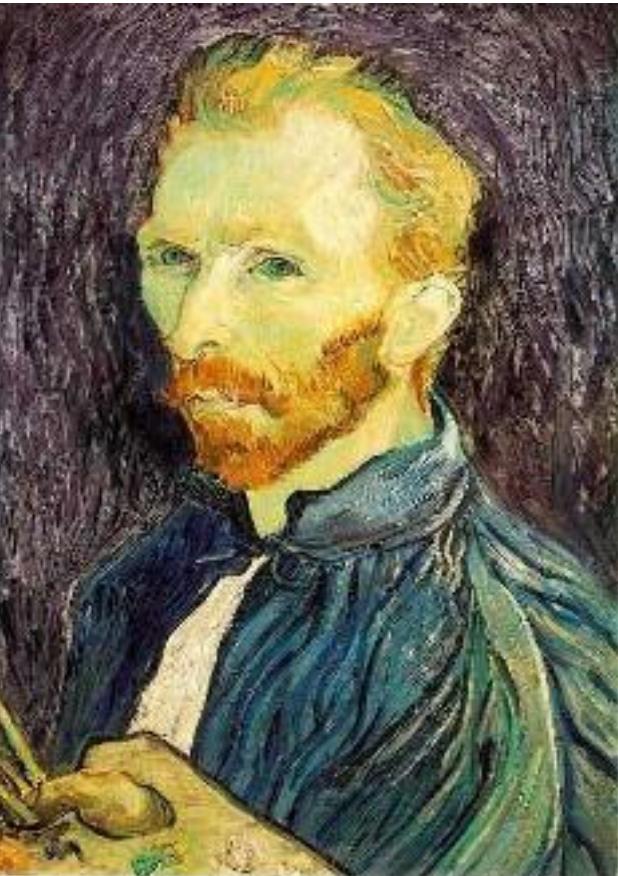
64x64

32x32

16x16



# Subsampling without pre-filtering



1/2

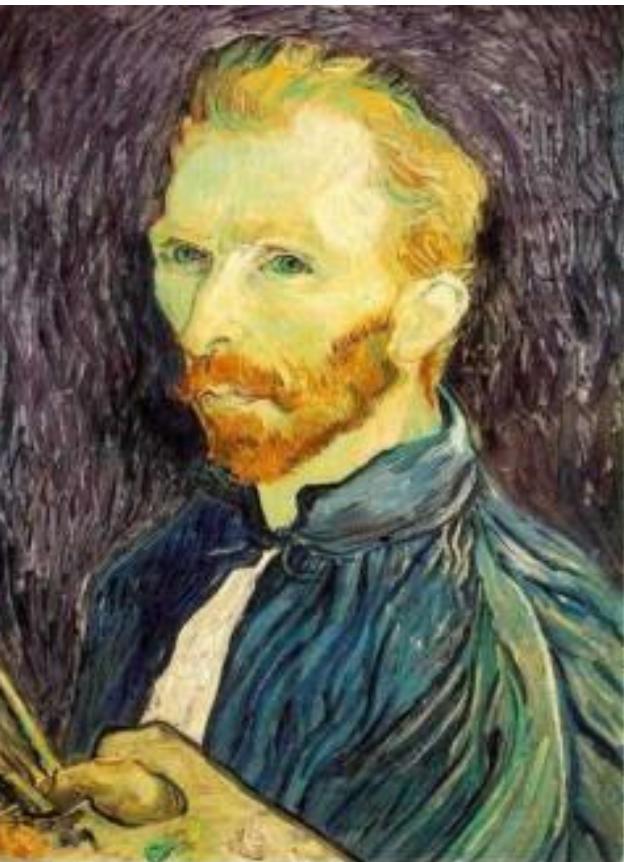


1/4 (2x zoom)



1/8 (4x zoom)

# Subsampling with Gaussian pre-filtering



Gaussian 1/2



G 1/4



G 1/8