



SurenManvelyan.com

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



Suren Manvelyan

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



Suren Manvelyan

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



Suren Manvelyan

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



SurenManvelyan.com

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



Suren Manvelyan

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



SurenManvelyan.com

By Suren Manvelyan, <http://www.surenmanvelyan.com/gallery/7116>



Indy



Indy

Heterochromia iridum

From Wikipedia, the free encyclopedia

Not to be confused with [Heterochromatin](#) or [Dichromatic](#) (disambiguation).

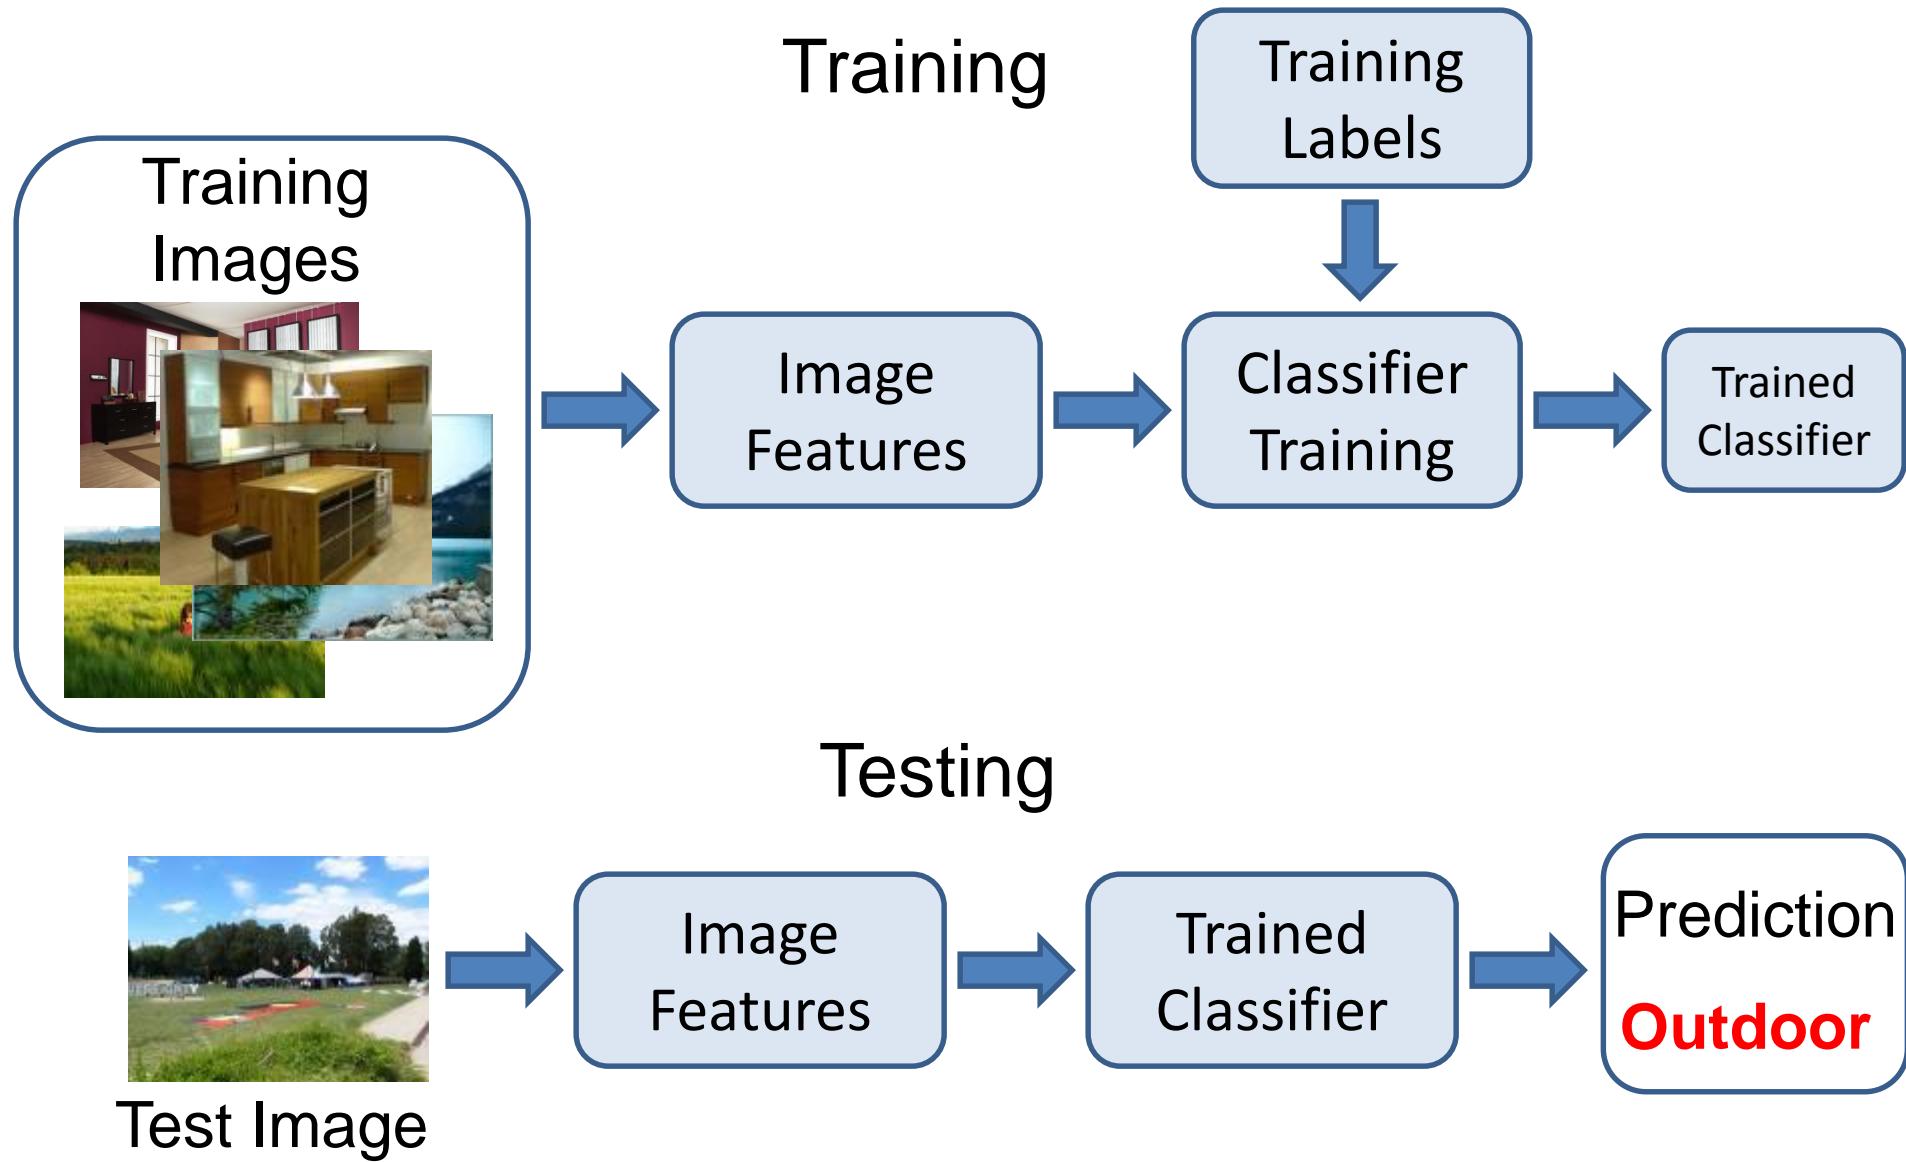
In anatomy, **heterochromia** ([ancient Greek](#): ἔτερος, héteros, different + χρώμα, chróma, color^[1]) is a difference in [coloration](#), usually of the [iris](#) but also of [hair](#) or [skin](#).

Heterochromia is a result of the relative excess or lack of [melanin](#) (a [pigment](#)). It may be [inherited](#), or caused by [genetic mosaicism](#), [chimerism](#), [disease](#), or [injury](#).^[2]

Heterochromia of the [eye](#) (***heterochromia iridis*** or ***heterochromia iridum***) is of three kinds. In *complete heterochromia*, one iris is a different color from the other. In *sectoral heterochromia*, part of one iris is a different color from its remainder and finally in "central heterochromia" there are spikes of different colours radiating from the pupil.

Heterochromia	
	
Complete heterochromia in human eyes: one brown and one green/hazel	
Classification and external resources	
Specialty	ophthalmology
ICD-10	Q13.2 ↗, H20.8 ↗, L67.1 ↗
ICD-9-CM	364.53 ↗
OMIM	142500 ↗
DiseasesDB	31289 ↗

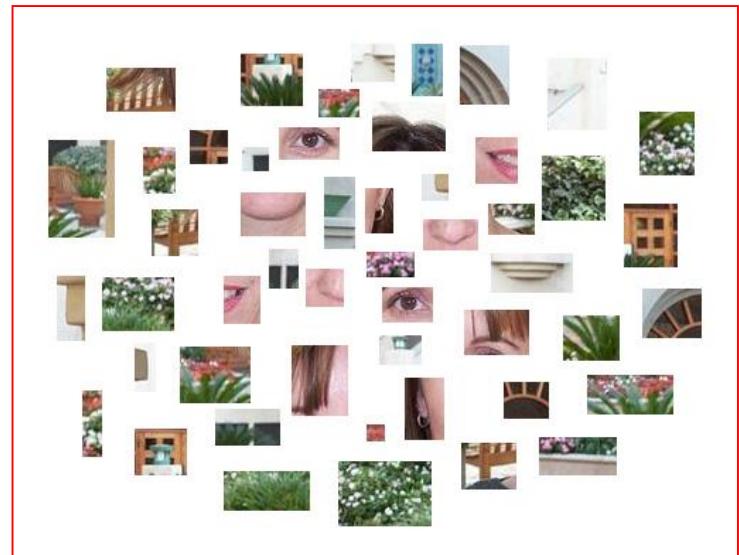
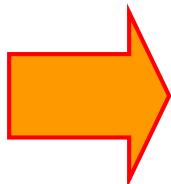
Image Categorization



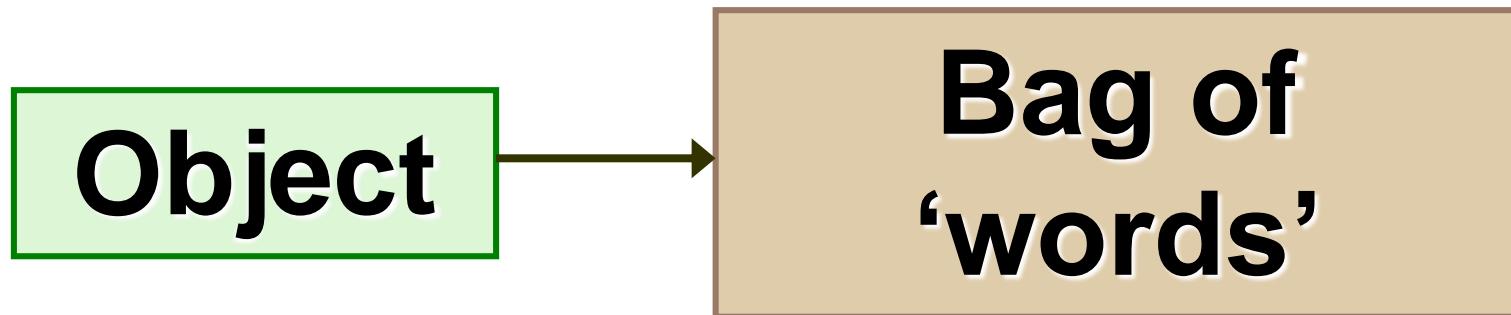
History of ideas in recognition

- 1960s – early 1990s: the geometric era
- 1990s: appearance-based models
- Mid-1990s: sliding window approaches
- Late 1990s: local features
- Early 2000s: parts-and-shape models
- Mid-2000s: bags of features

Bag-of-features models

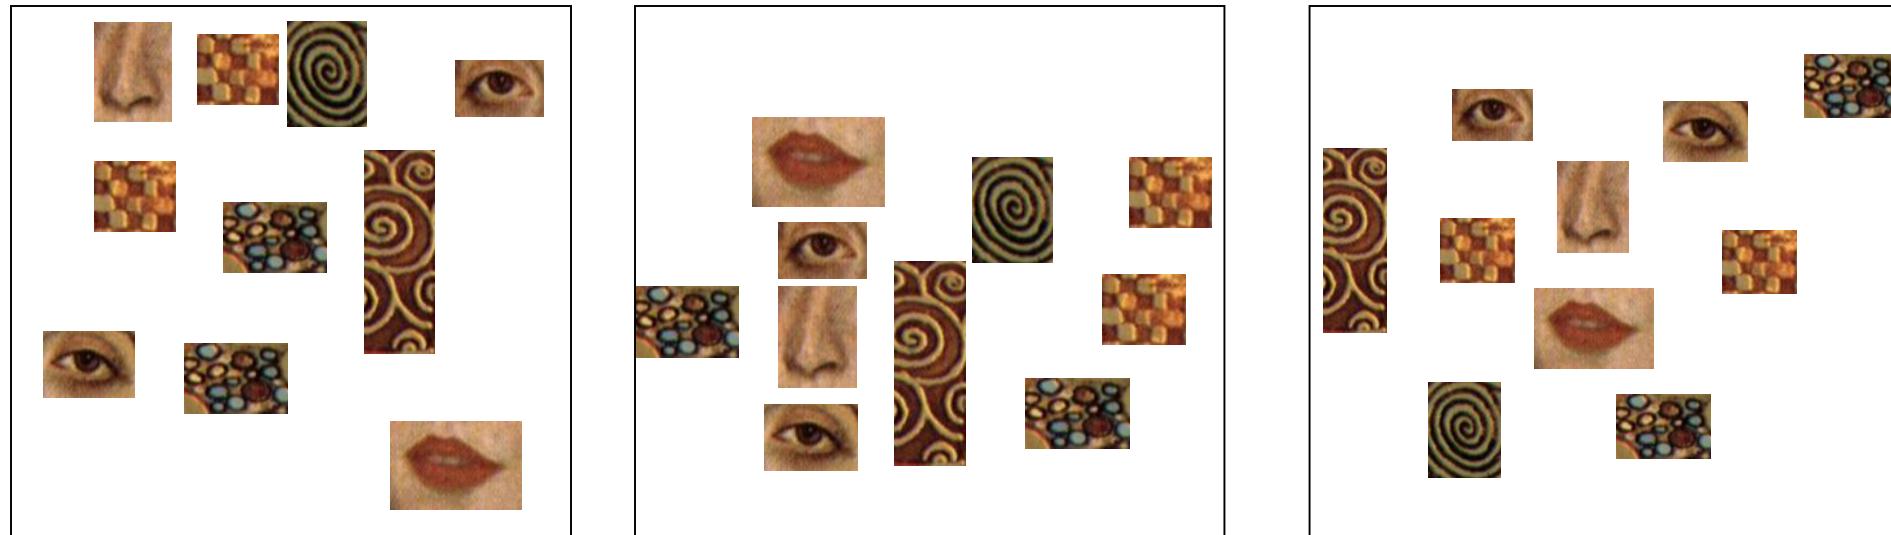


Bag-of-features models



Objects as texture

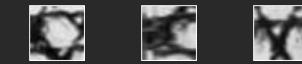
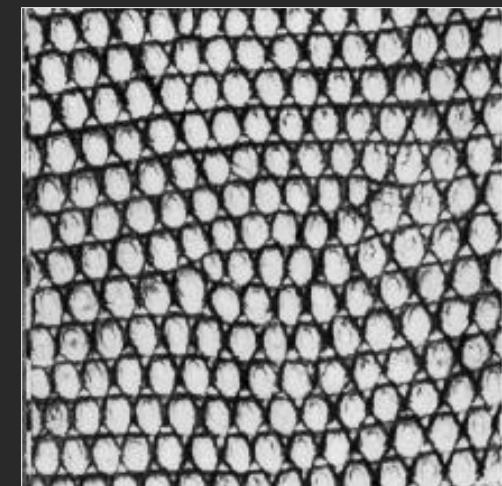
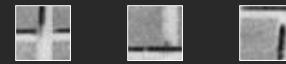
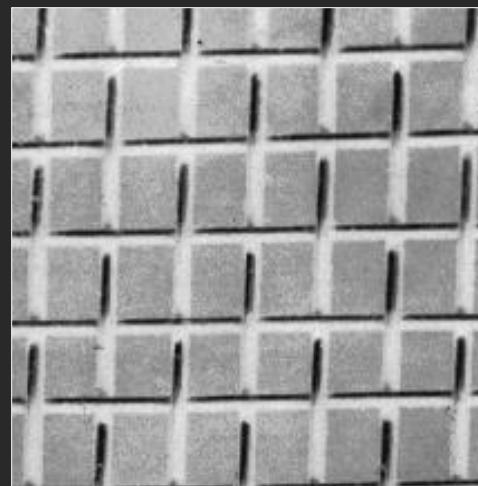
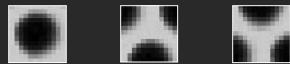
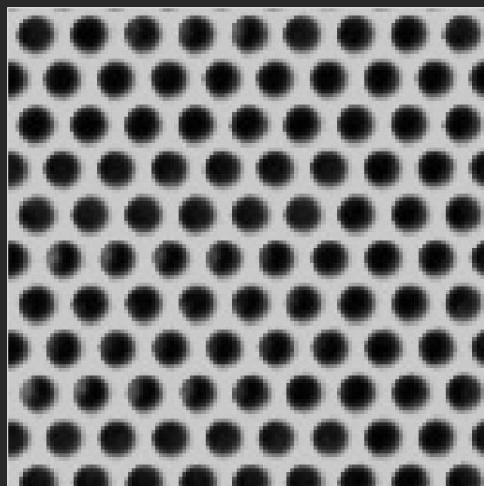
- All of these are treated as being the same



- No distinction between foreground and background: scene recognition?

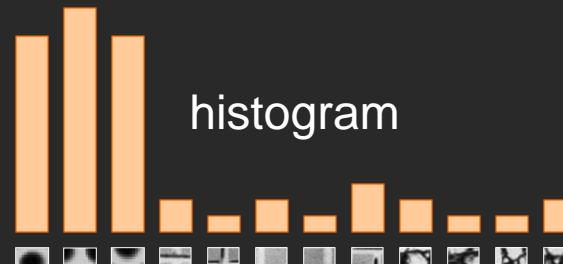
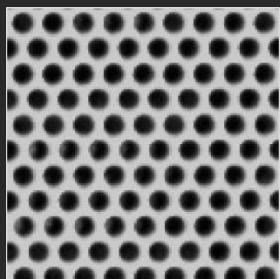
Origin 1: Texture recognition

- Texture is characterized by the repetition of basic elements or *textons*
- For stochastic textures, it is the identity of the textons, not their spatial arrangement, that matters

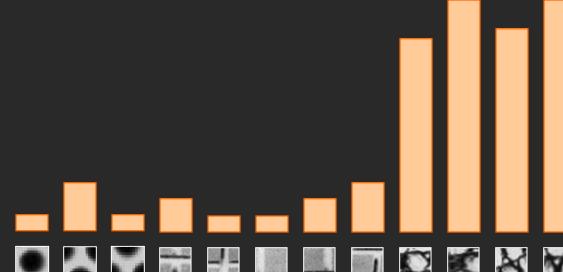
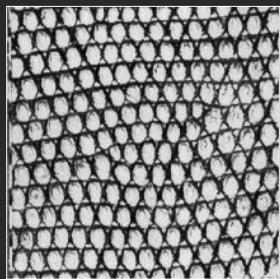
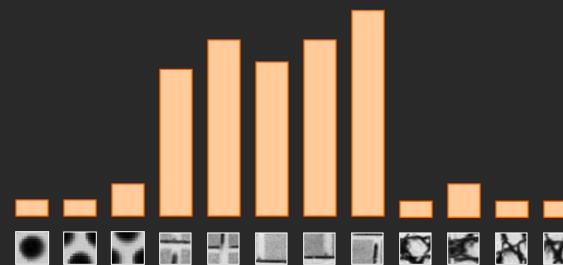
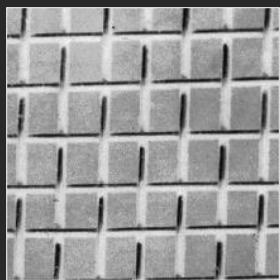


Julesz, 1981; Cula & Dana, 2001; Leung & Malik 2001; Mori, Belongie & Malik, 2001; Schmid 2001; Varma & Zisserman, 2002, 2003; Lazebnik, Schmid & Ponce, 2003

Origin 1: Texture recognition



Universal texton dictionary



Julesz, 1981; Cula & Dana, 2001; Leung & Malik 2001; Mori, Belongie & Malik, 2001; Schmid 2001; Varma & Zisserman, 2002, 2003; Lazebnik, Schmid & Ponce, 2003

Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)

Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)



Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)



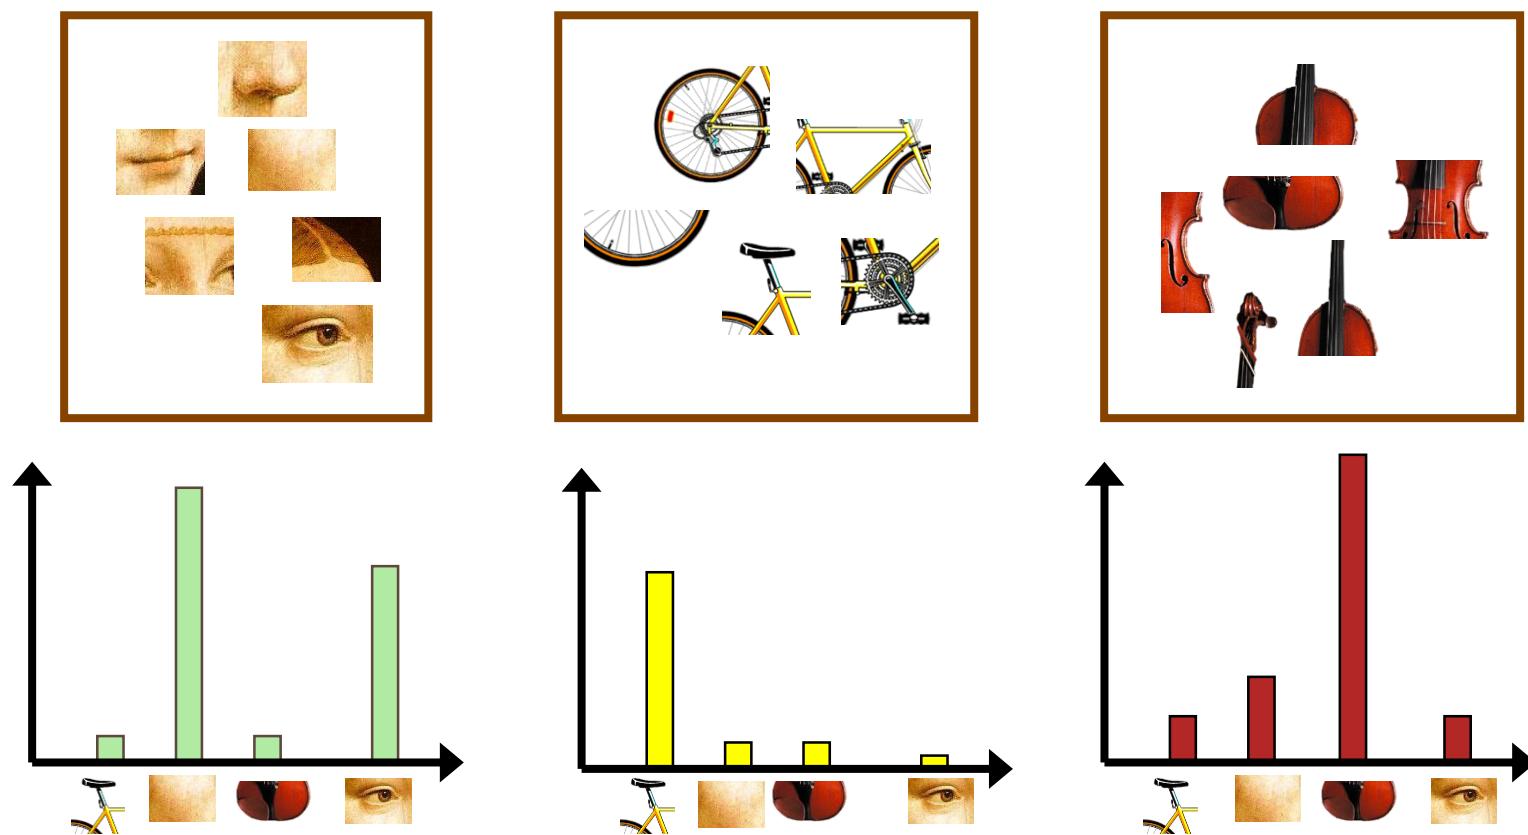
Origin 2: Bag-of-words models

- Orderless document representation: frequencies of words from a dictionary Salton & McGill (1983)



Bag-of-features steps

1. Extract features
2. Learn “visual vocabulary”
3. Quantize features using visual vocabulary
4. Represent images by frequencies of “visual words”

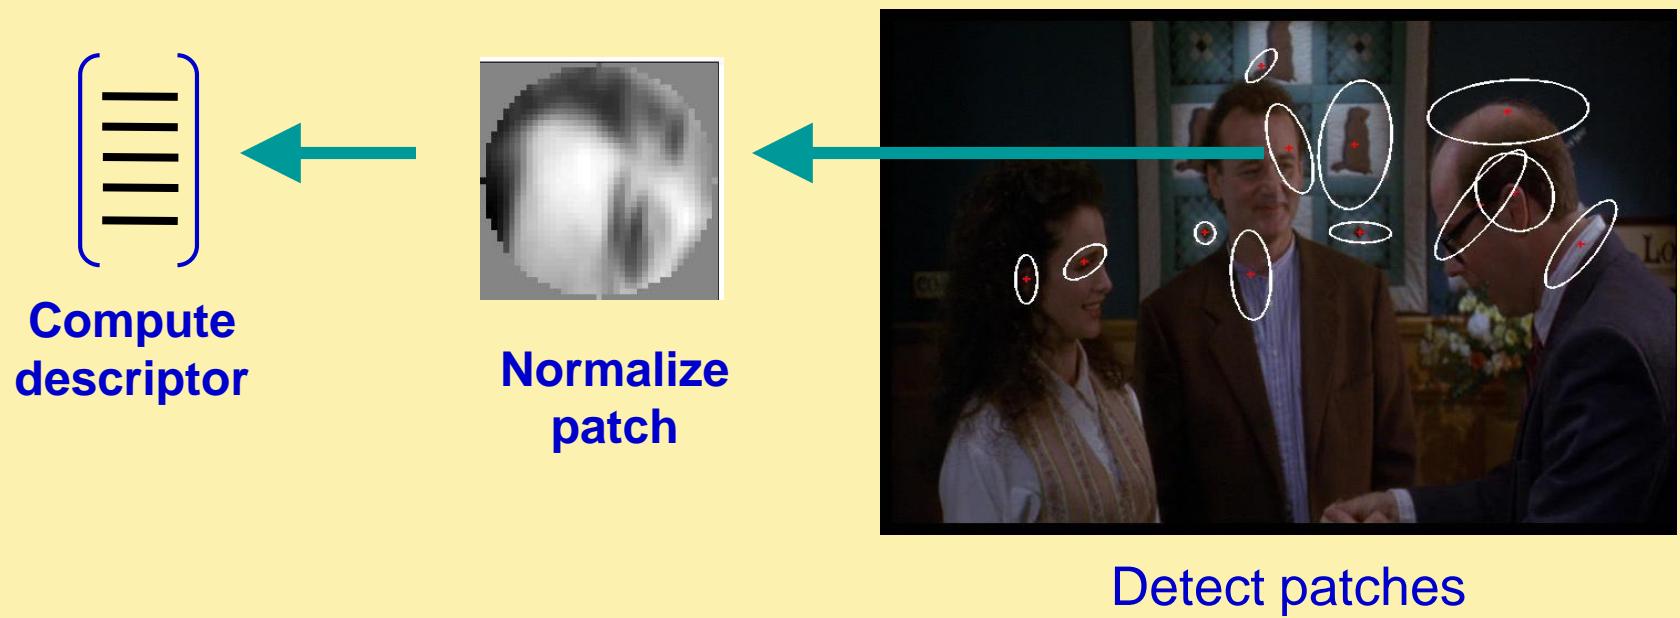


1. Feature extraction

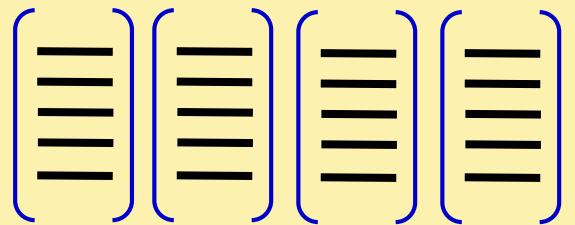
- Regular grid or interest regions



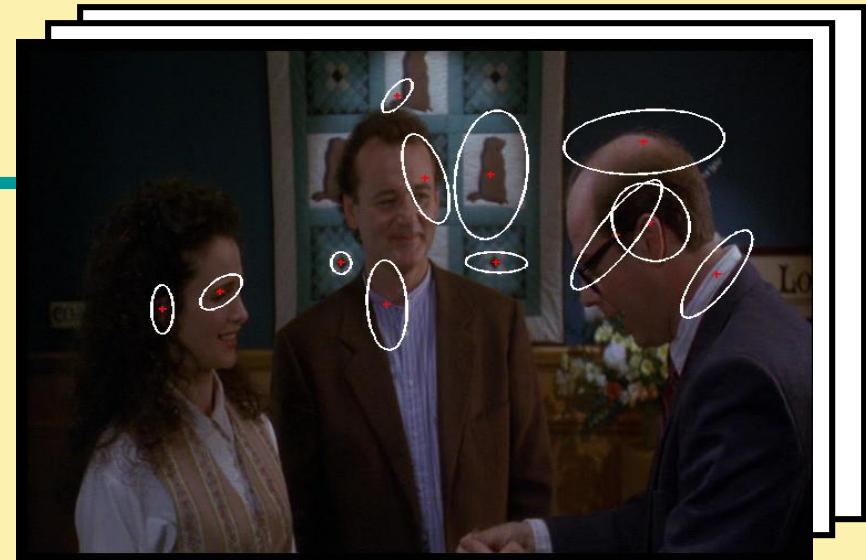
1. Feature extraction



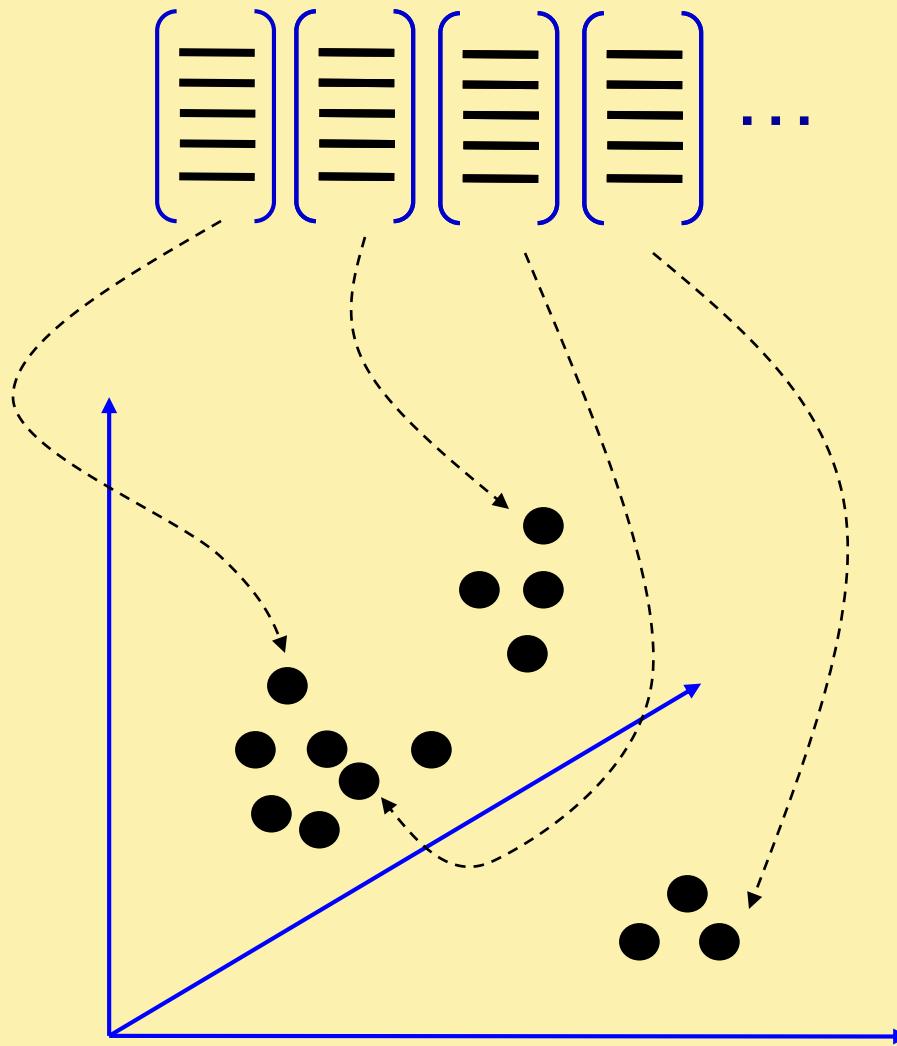
1. Feature extraction



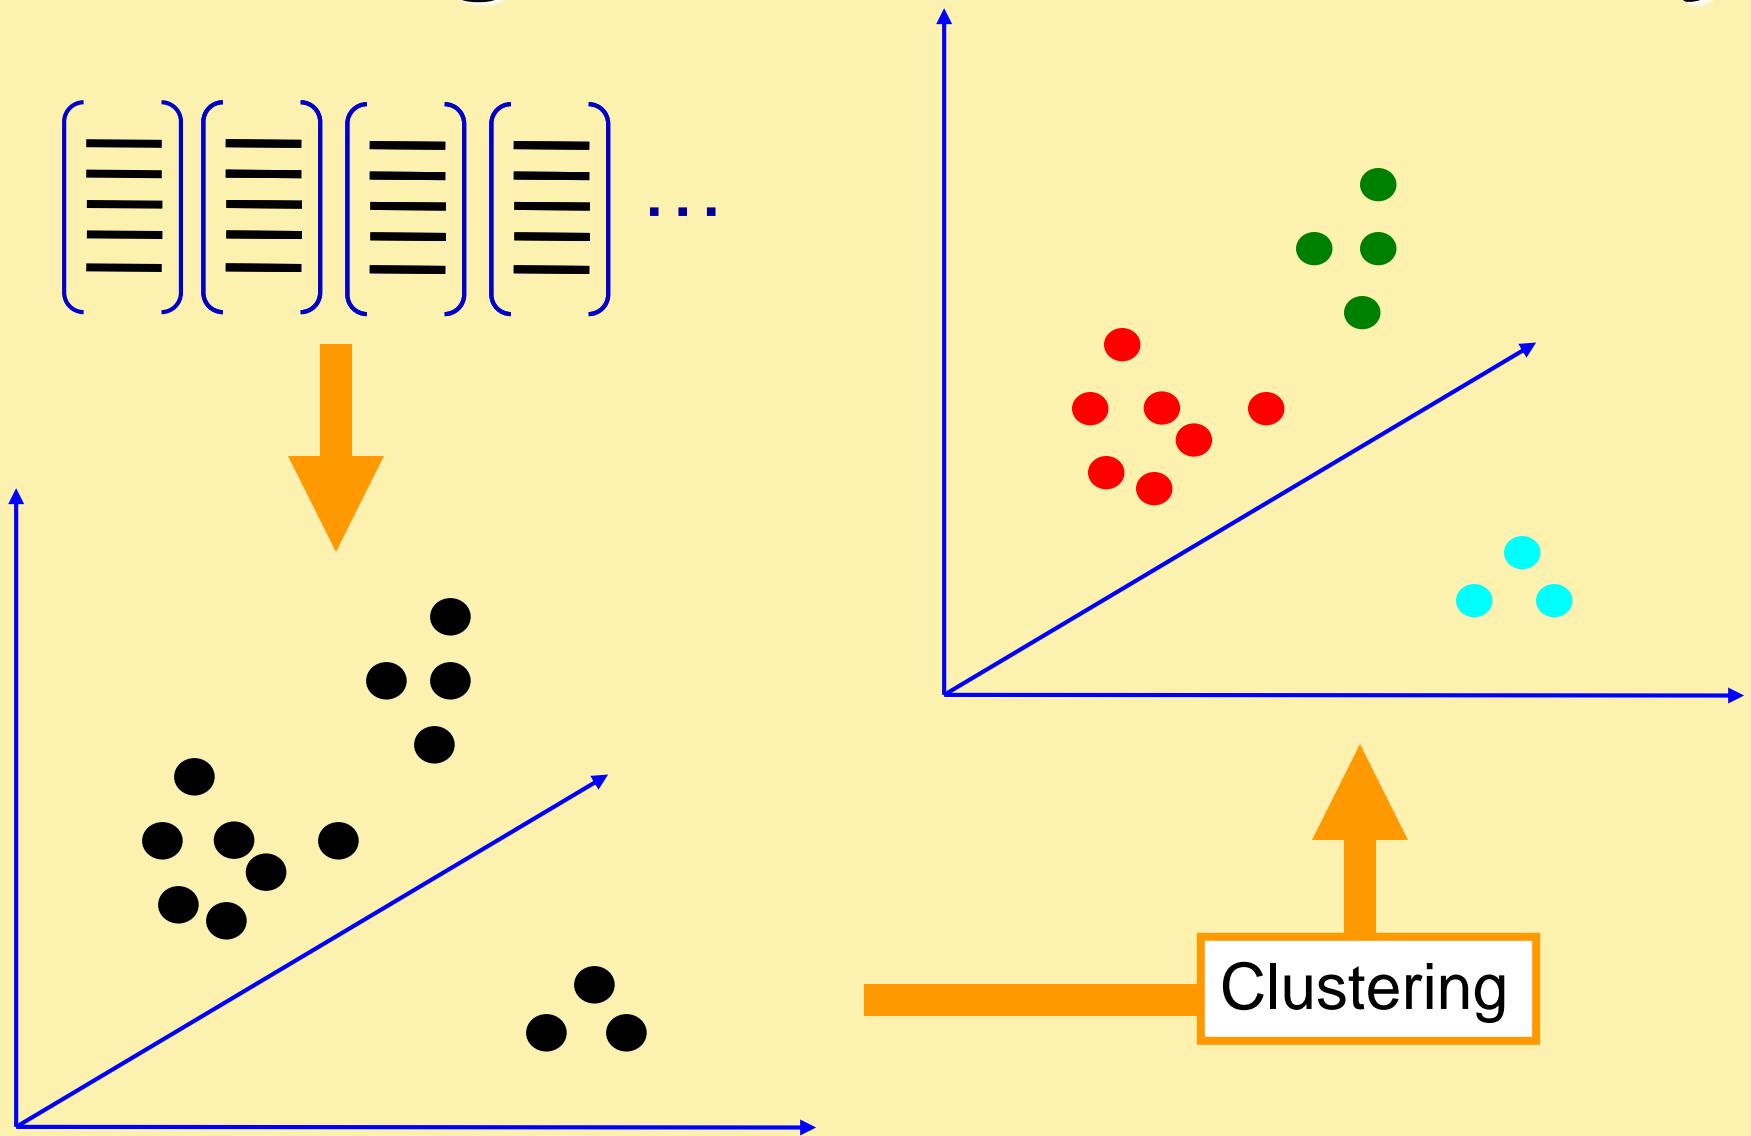
... ←



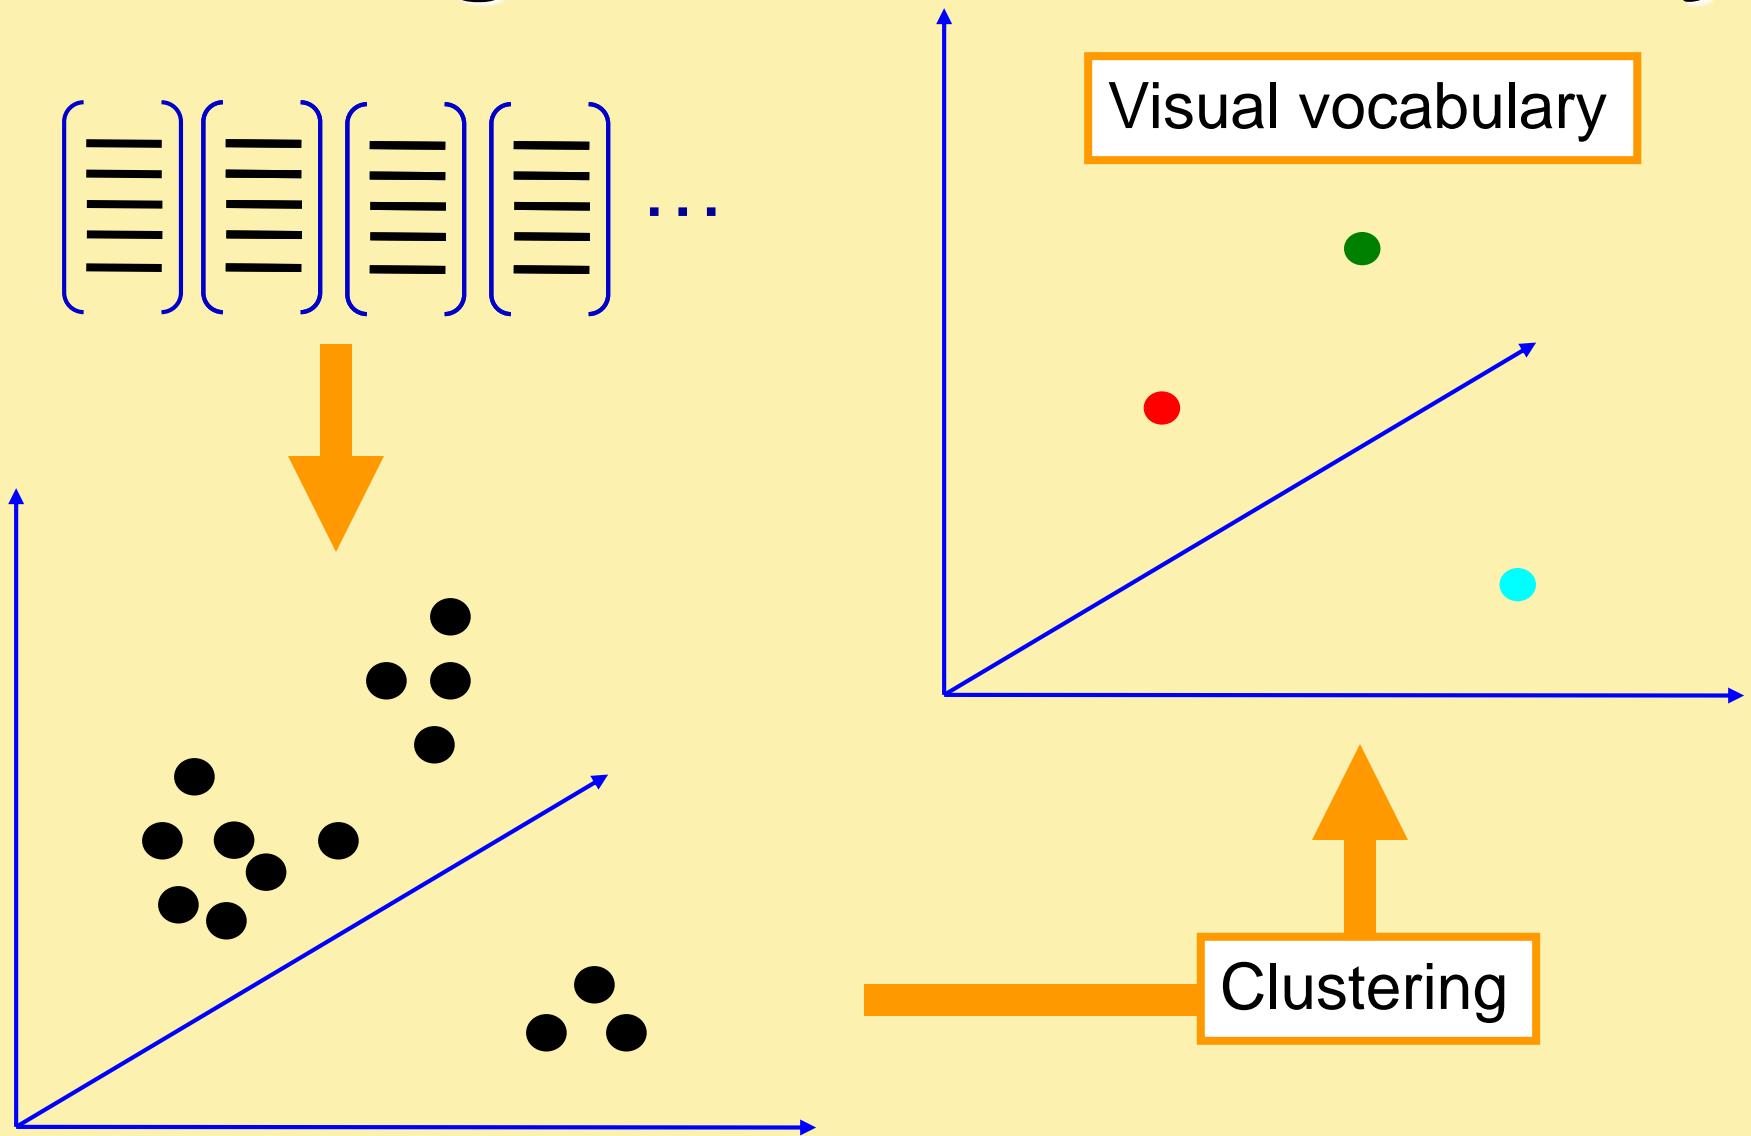
2. Learning the visual vocabulary



2. Learning the visual vocabulary



2. Learning the visual vocabulary



K-means clustering

- Want to minimize sum of squared Euclidean distances between points x_i and their nearest cluster centers m_k

$$D(X, M) = \sum_{\text{cluster } k} \sum_{\substack{\text{point } i \text{ in} \\ \text{cluster } k}} (x_i - m_k)^2$$

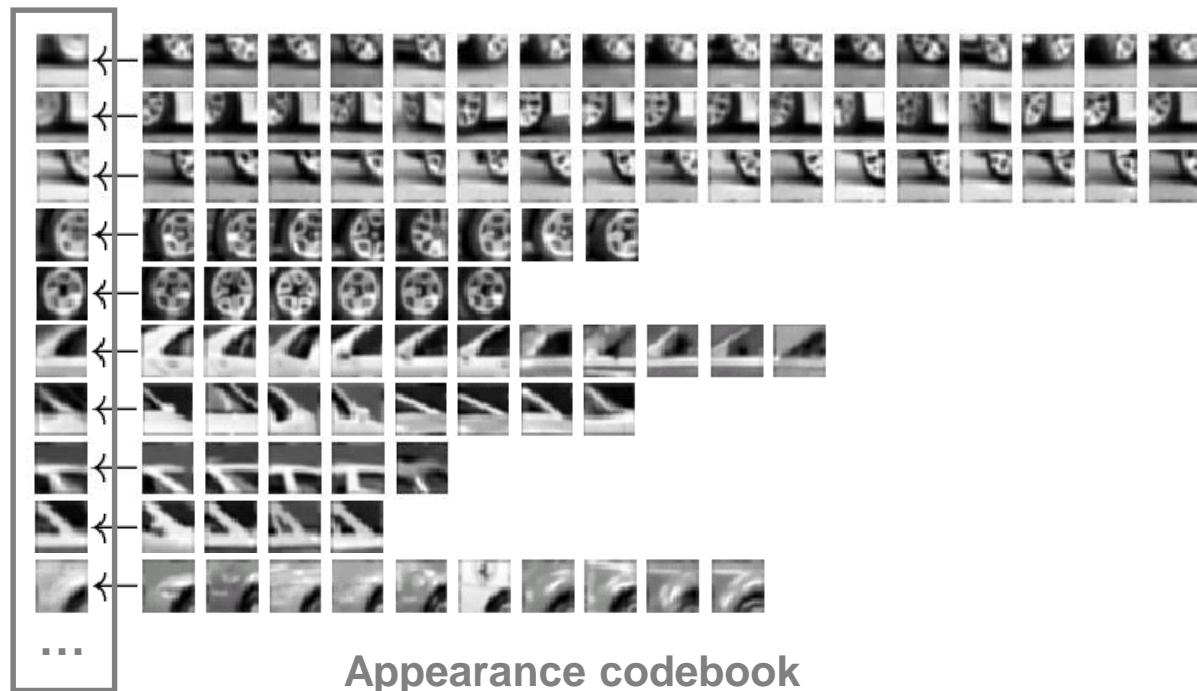
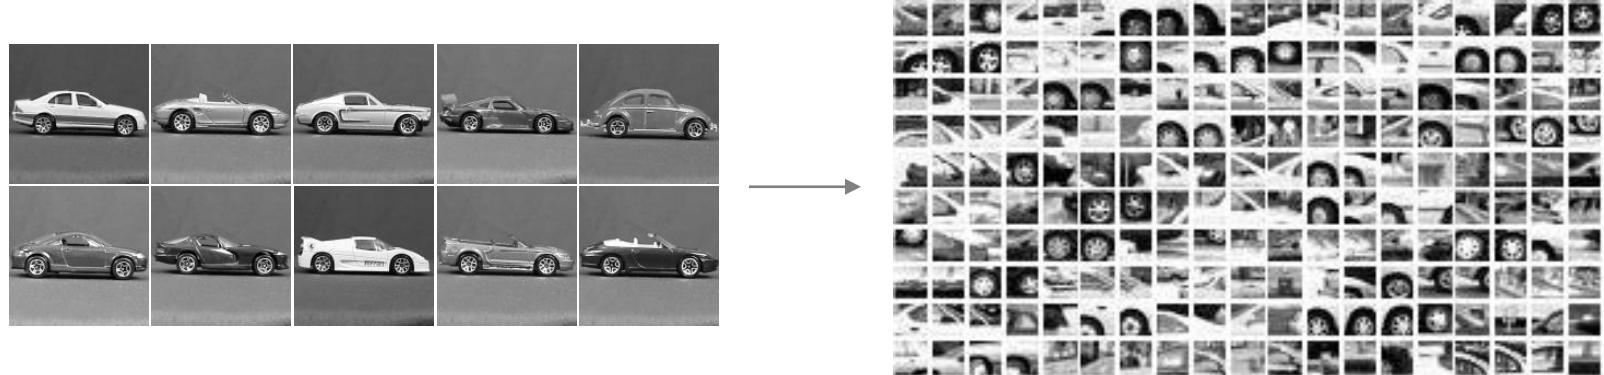
Algorithm:

- Randomly initialize K cluster centers
- Iterate until convergence:
 - Assign each data point to the nearest center
 - Recompute each cluster center as the mean of all points assigned to it

Clustering and vector quantization

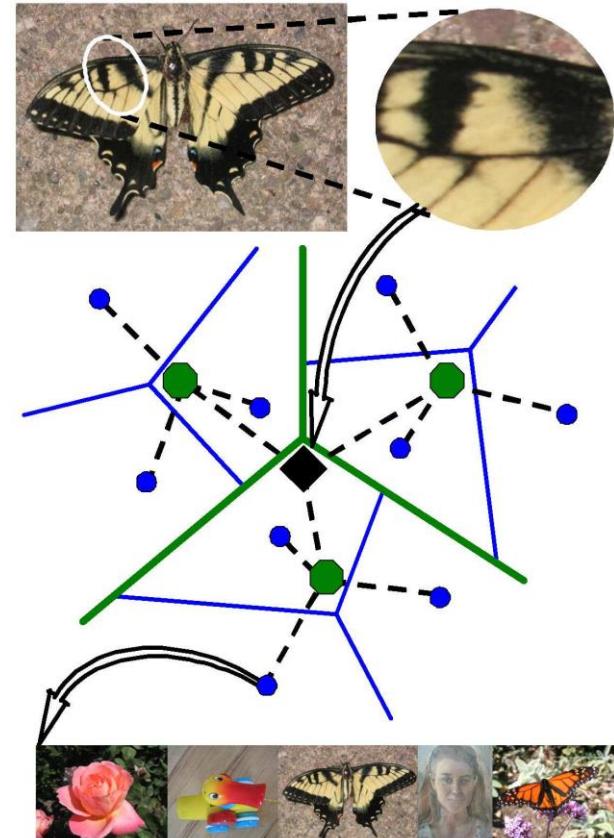
- Clustering is a common method for learning a visual vocabulary or codebook
 - Unsupervised learning process
 - Each cluster center produced by k-means becomes a codevector
 - Codebook can be learned on separate training set
 - Provided the training set is sufficiently representative, the codebook will be “universal”
- The codebook is used for quantizing features
 - A *vector quantizer* takes a feature vector and maps it to the index of the nearest codevector in a codebook
 - Codebook = visual vocabulary
 - Codevector = visual word

Example codebook

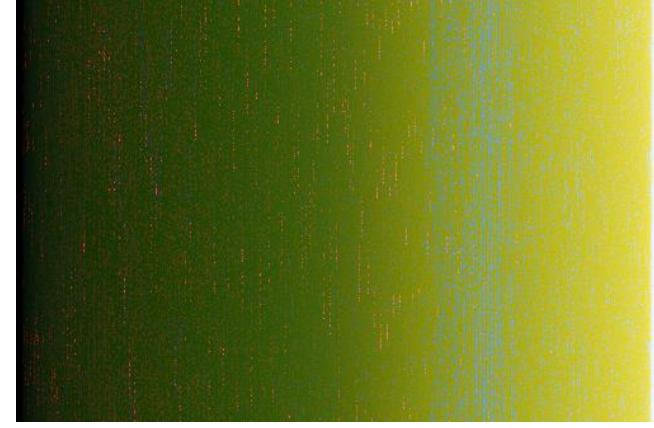
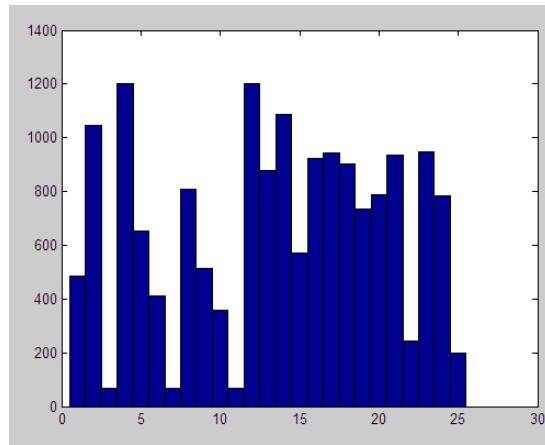


Visual vocabularies: Issues

- How to choose vocabulary size?
 - Too small: visual words not representative of all patches
 - Too large: quantization artifacts, overfitting
- Computational efficiency
 - Vocabulary trees
(Nister & Stewenius, 2006)



But what about layout?



All of these images have the same color histogram

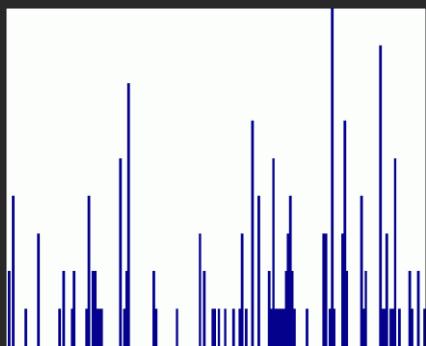
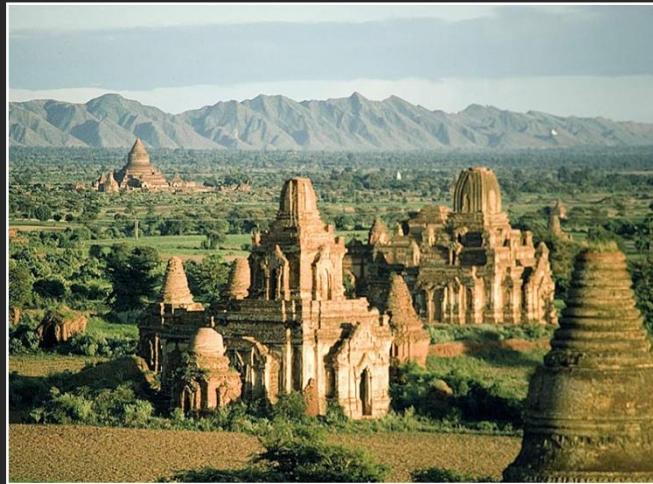
Spatial pyramid



Compute histogram in each spatial bin

Spatial pyramid representation

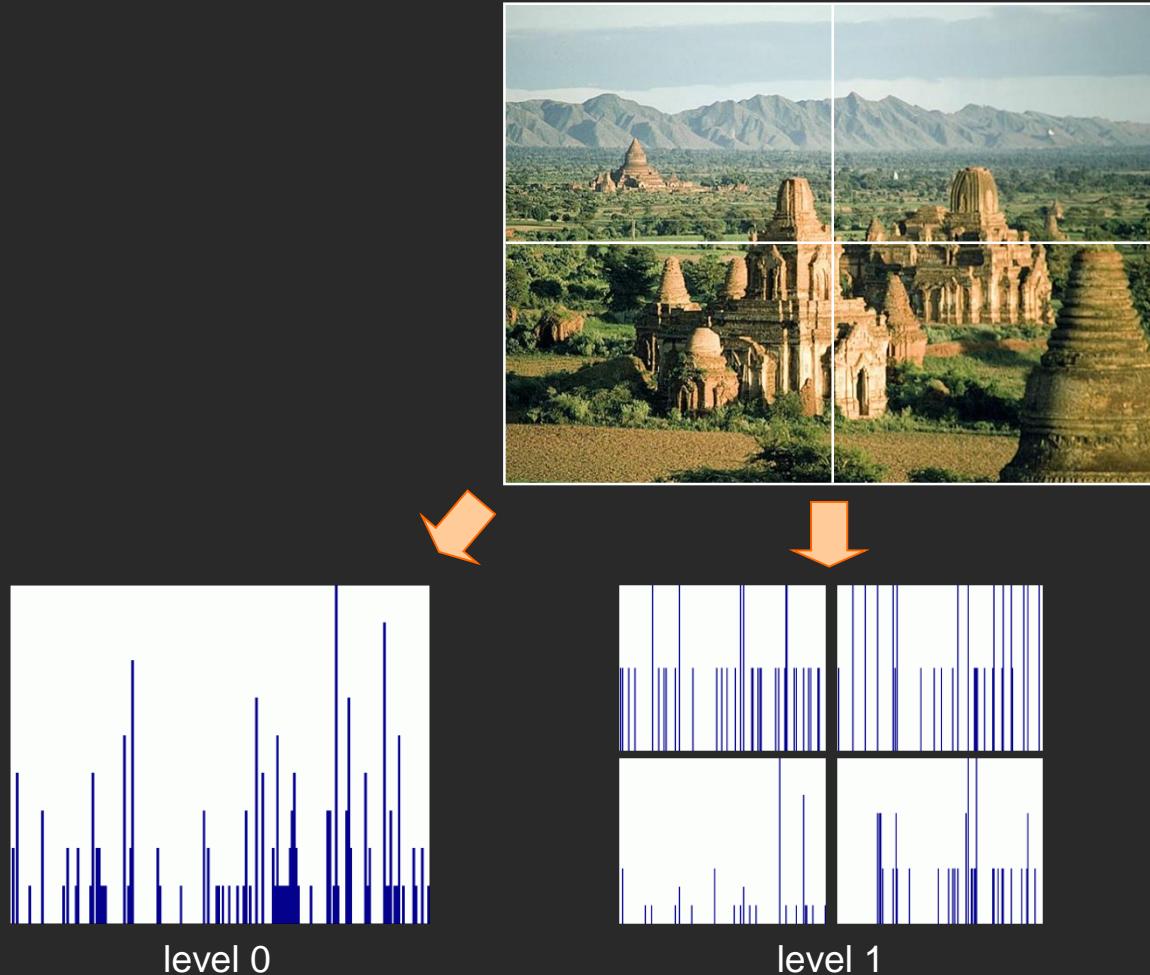
- Extension of a bag of features
- Locally orderless representation at several levels of resolution



level 0

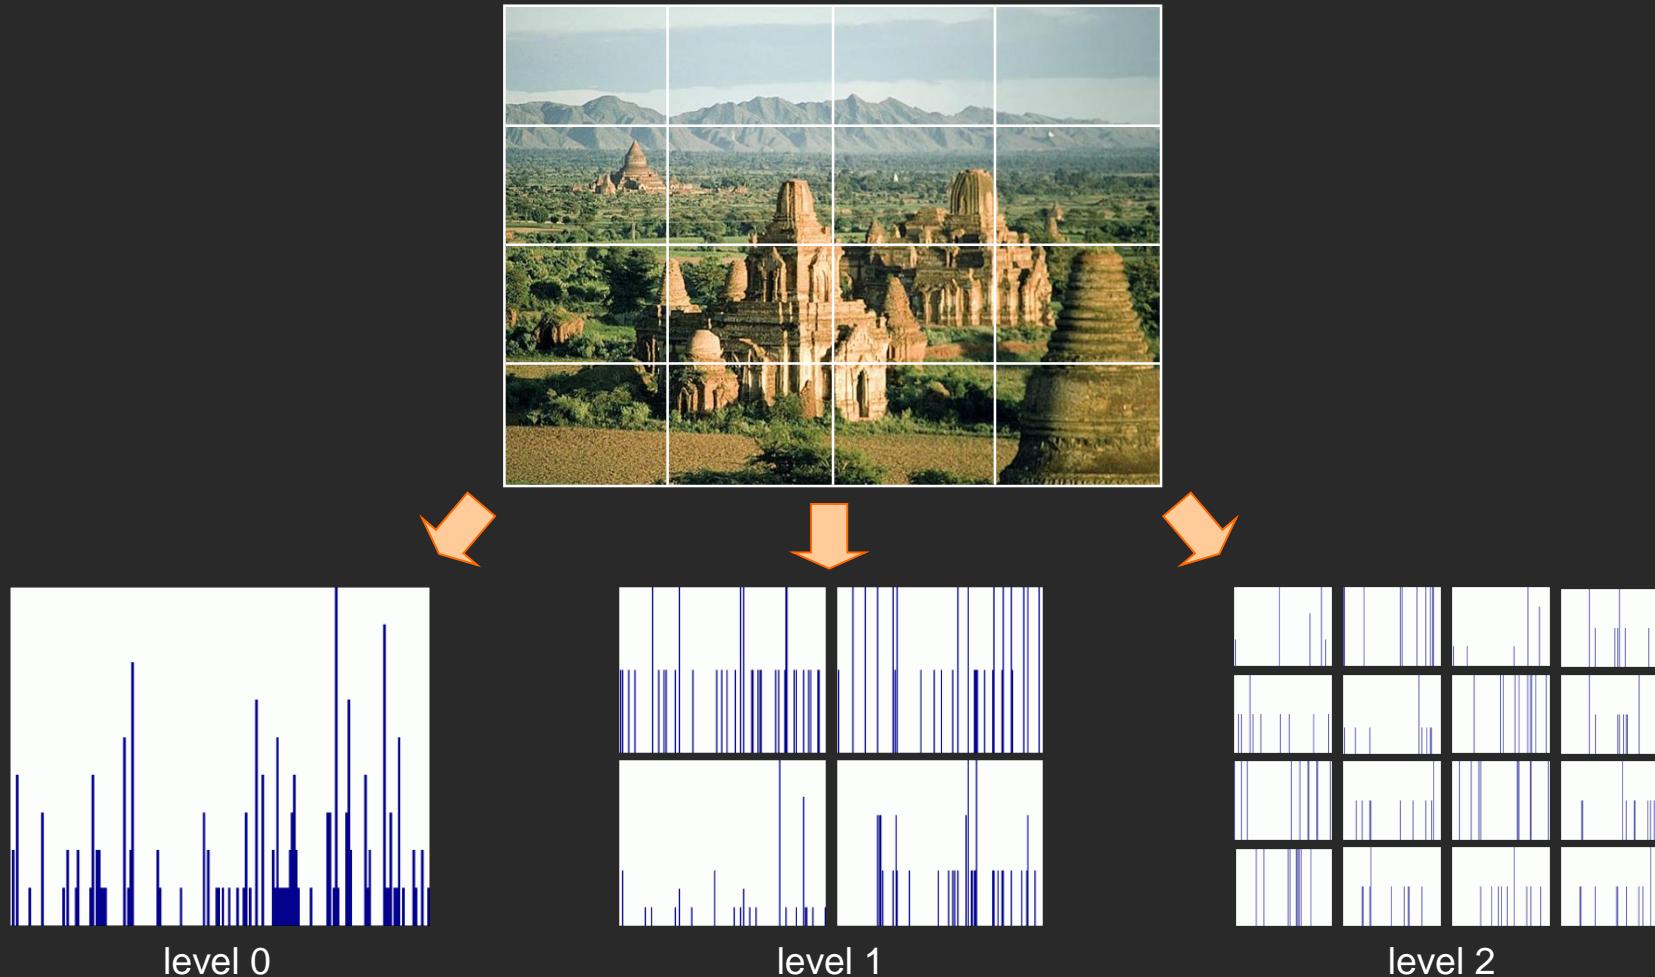
Spatial pyramid representation

- Extension of a bag of features
- Locally orderless representation at several levels of resolution

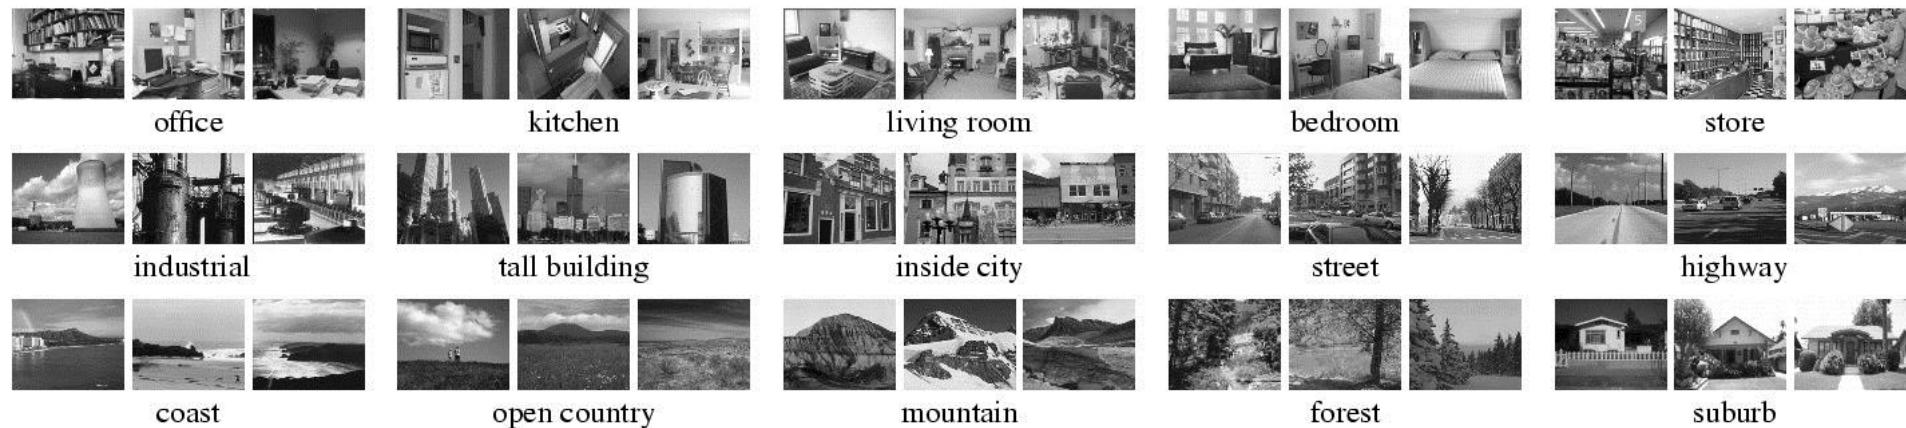


Spatial pyramid representation

- Extension of a bag of features
- Locally orderless representation at several levels of resolution



Scene category dataset



Multi-class classification results
(100 training images per class)

Level	Weak features (vocabulary size: 16)		Strong features (vocabulary size: 200)	
	Single-level	Pyramid	Single-level	Pyramid
0 (1×1)	45.3 ± 0.5		72.2 ± 0.6	
1 (2×2)	53.6 ± 0.3	56.2 ± 0.6	77.9 ± 0.6	79.0 ± 0.5
2 (4×4)	61.7 ± 0.6	64.7 ± 0.7	79.4 ± 0.3	81.1 ± 0.3
3 (8×8)	63.3 ± 0.8	66.8 ± 0.6	77.2 ± 0.4	80.7 ± 0.3

Caltech101 dataset

http://www.vision.caltech.edu/Image_Datasets/Caltech101/Caltech101.html

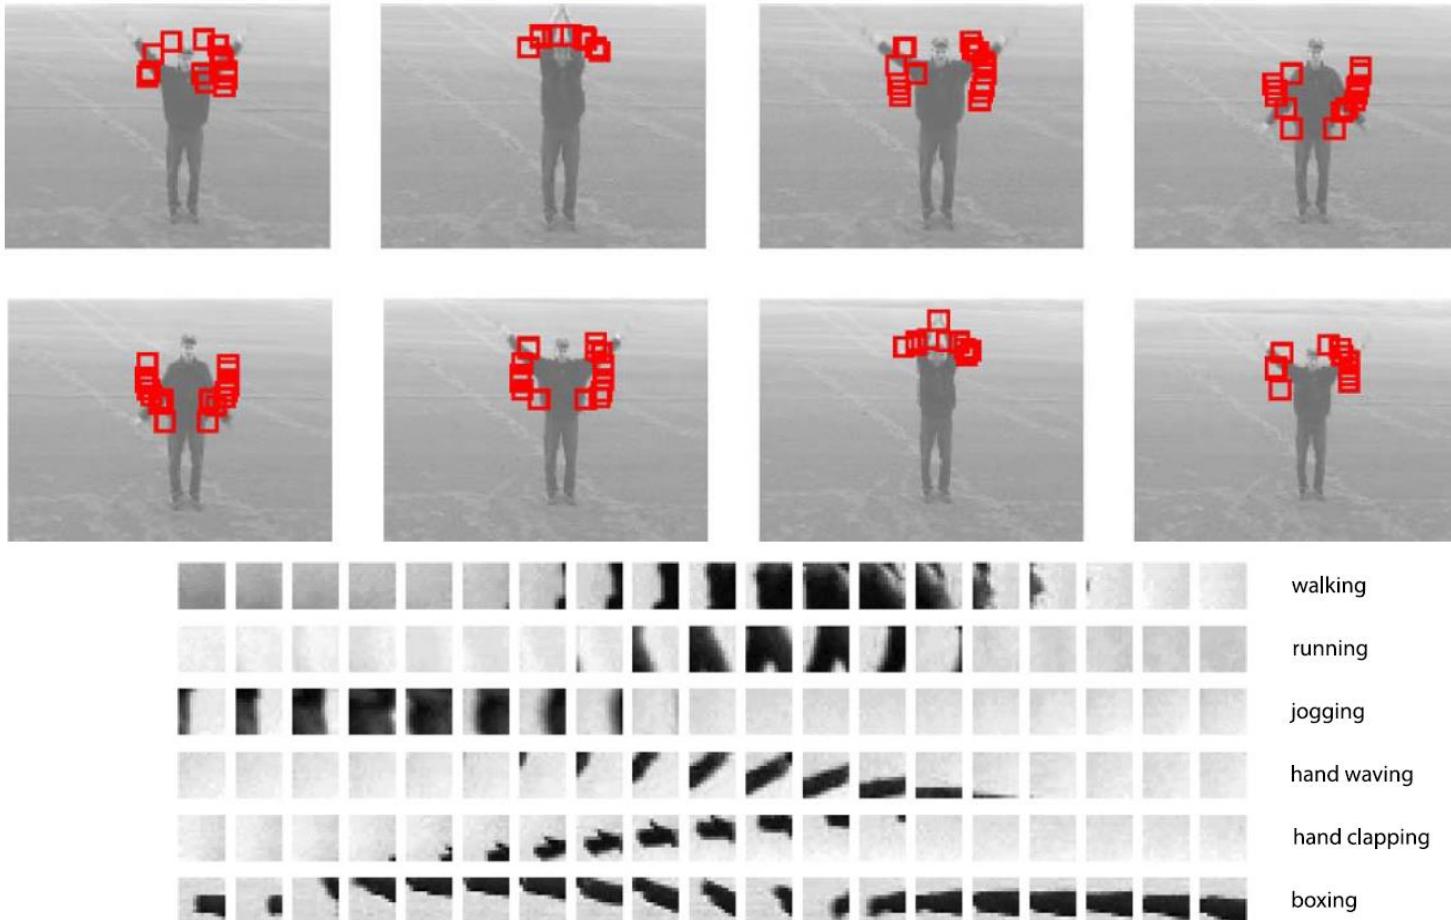


Multi-class classification results (30 training images per class)

	Weak features (16)		Strong features (200)	
Level	Single-level	Pyramid	Single-level	Pyramid
0	15.5 ± 0.9		41.2 ± 1.2	
1	31.4 ± 1.2	32.8 ± 1.3	55.9 ± 0.9	57.0 ± 0.8
2	47.2 ± 1.1	49.3 ± 1.4	63.6 ± 0.9	64.6 ± 0.8
3	52.2 ± 0.8	54.0 ± 1.1	60.3 ± 0.9	64.6 ± 0.7

Bags of features for action recognition

Space-time interest points



Juan Carlos Niebles, Hongcheng Wang and Li Fei-Fei, [Unsupervised Learning of Human Action Categories Using Spatial-Temporal Words](#), IJCV 2008.

History of ideas in recognition

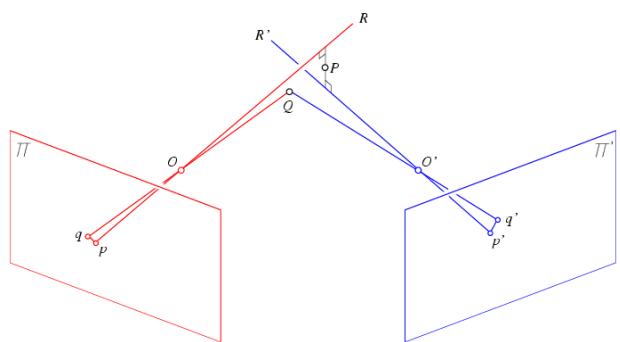
- 1960s – early 1990s: the geometric era
- 1990s: appearance-based models
- Mid-1990s: sliding window approaches
- Late 1990s: local features
- Early 2000s: parts-and-shape models
- Mid-2000s: bags of features
- Present trends: combination of local and global methods, context, *deep learning*

Large-scale Instance Retrieval

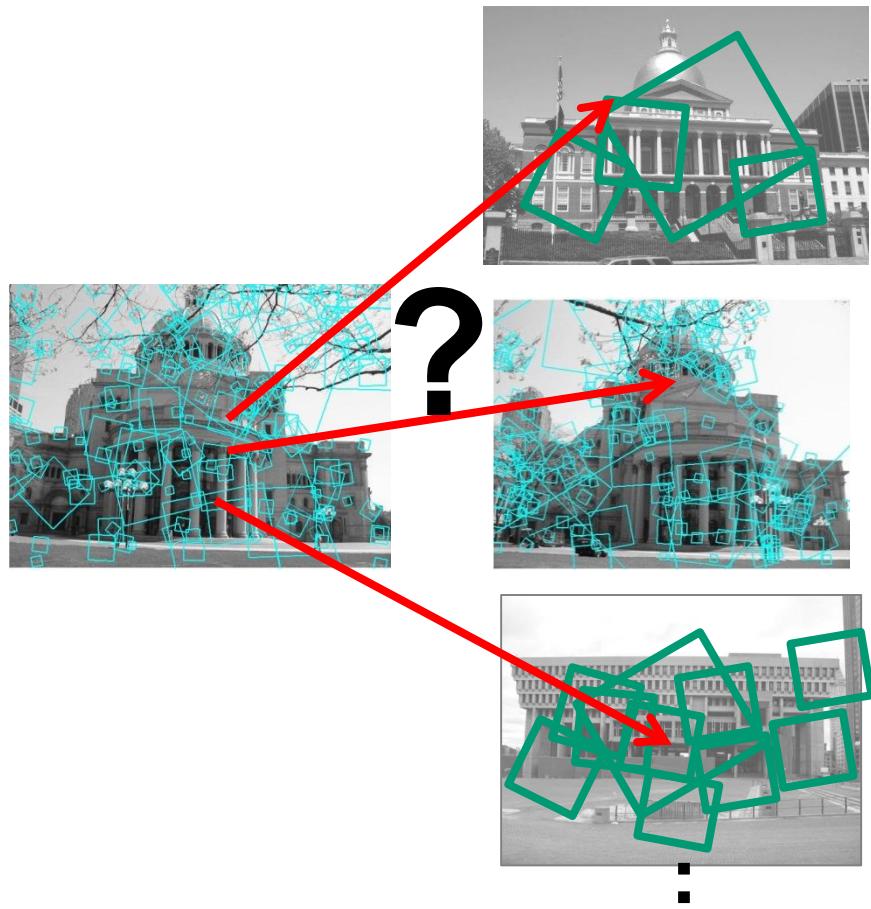
Computer Vision

James Hays

Multi-view matching



vs



Matching two given
views for depth

Search for a matching
view for recognition

Video Google System

1. Collect all words within query region
2. Inverted file index to find relevant frames
3. Compare word counts
4. Spatial verification

Sivic & Zisserman, ICCV 2003

- Demo online at :
<http://www.robots.ox.ac.uk/~vgg/research/vgoogle/index.html>



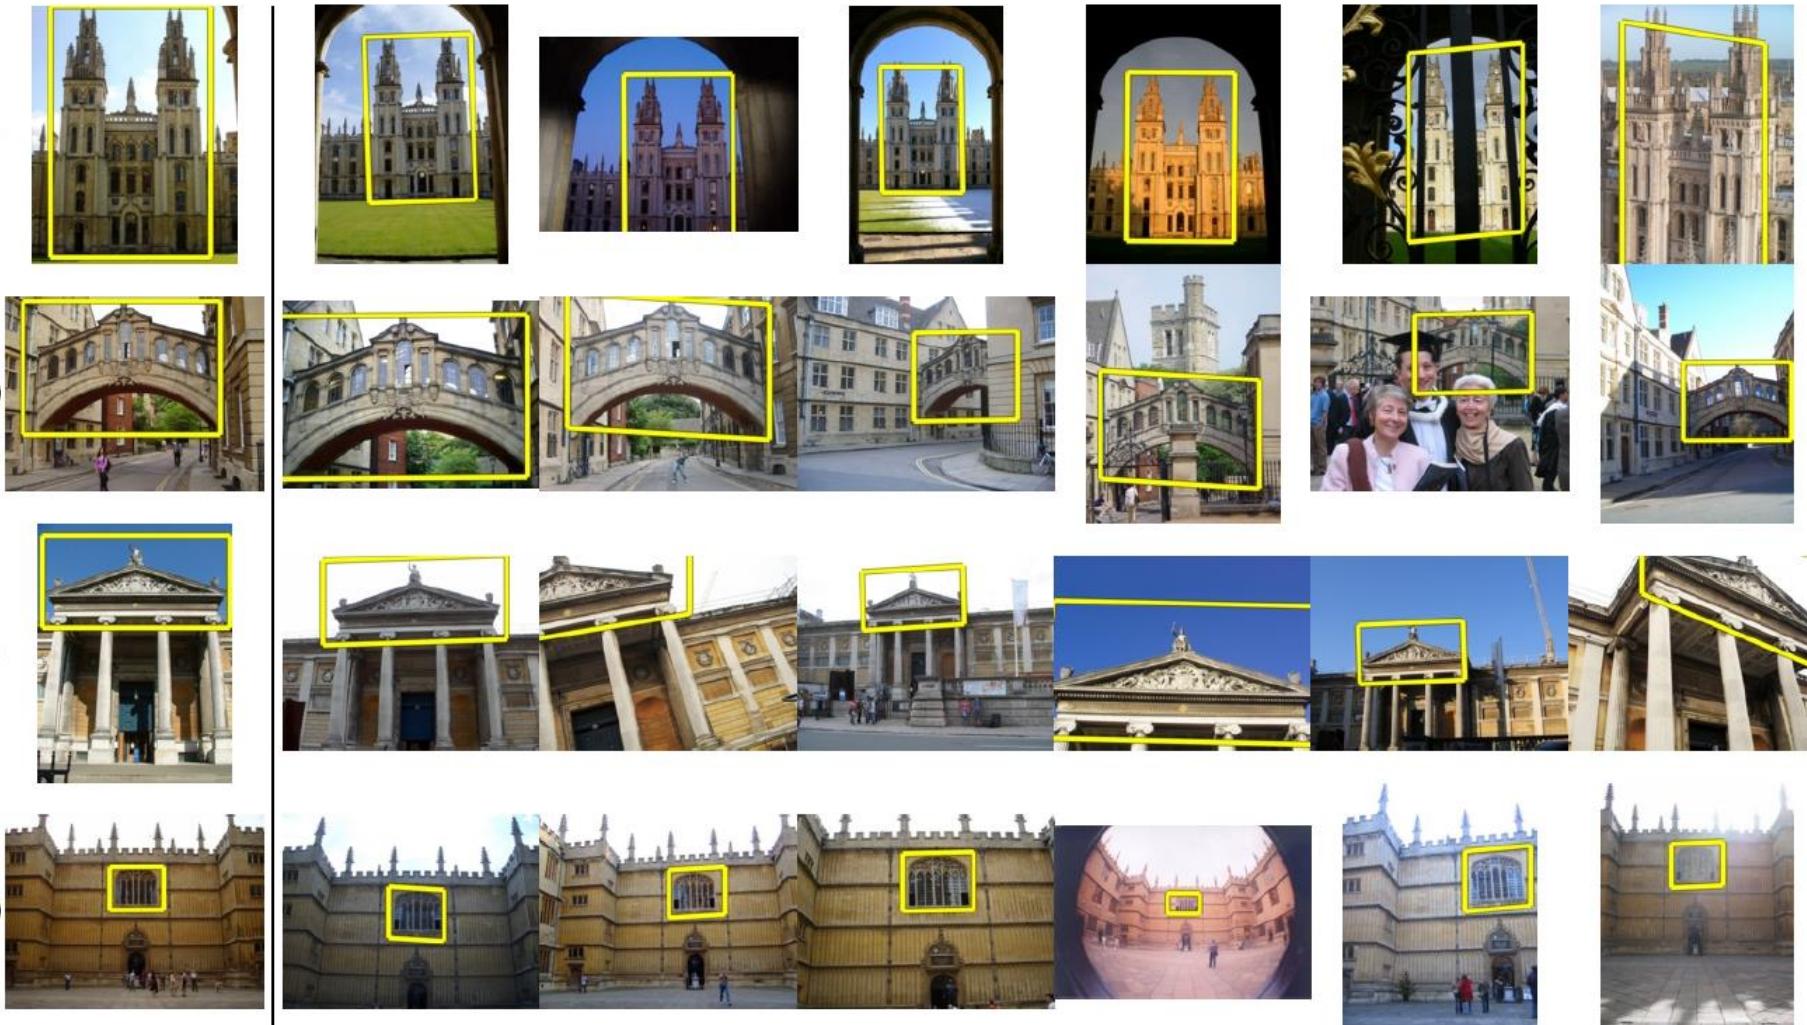
Query region



Kristen Grauman

Retrieved frames

Application: Large-Scale Retrieval



Query

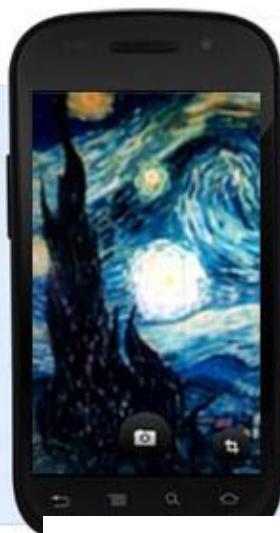
Results from 5k Flickr images (demo available for 100k set)



Google Goggles

Use pictures to search the web.

Watch a video



Get Google Goggles

Android (1.6+ required)

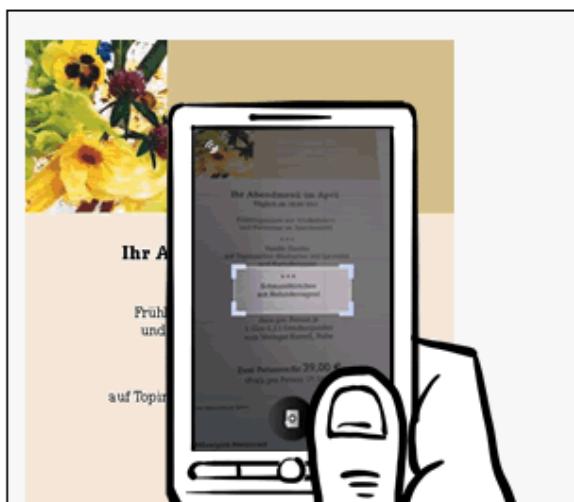
Download from [Android Market](#).

[Send Goggles to Android phone](#)

New! iPhone (iOS 4.0 required)

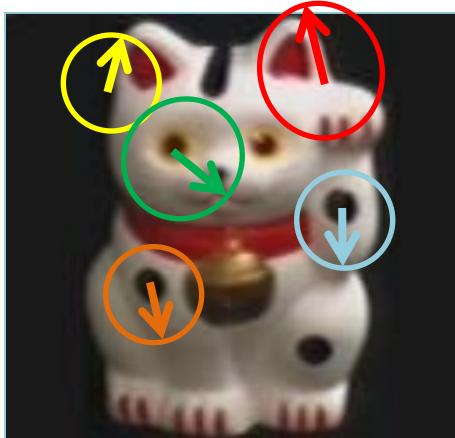
Download [from the App Store](#).

[Send Goggles to iPhone](#)

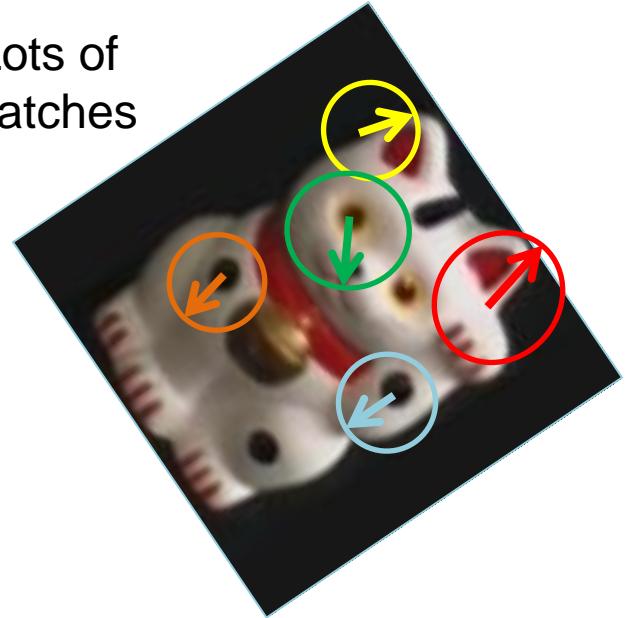


Simple idea

See how many keypoints
are close to keypoints in
each other image



Lots of
Matches



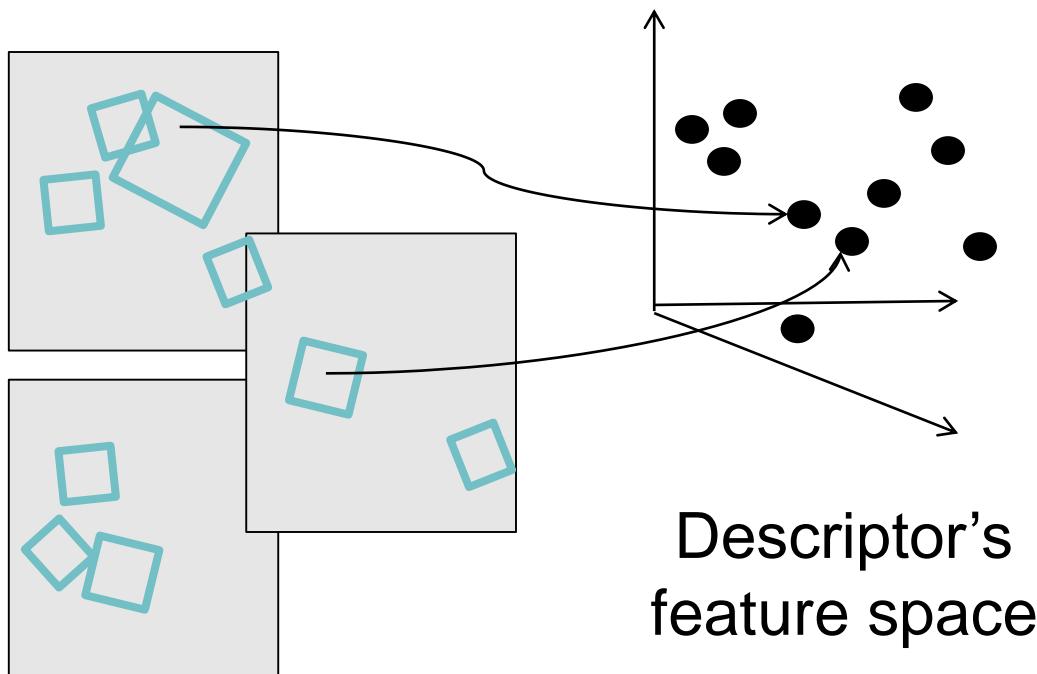
Few or No
Matches



But this will be really, really slow!

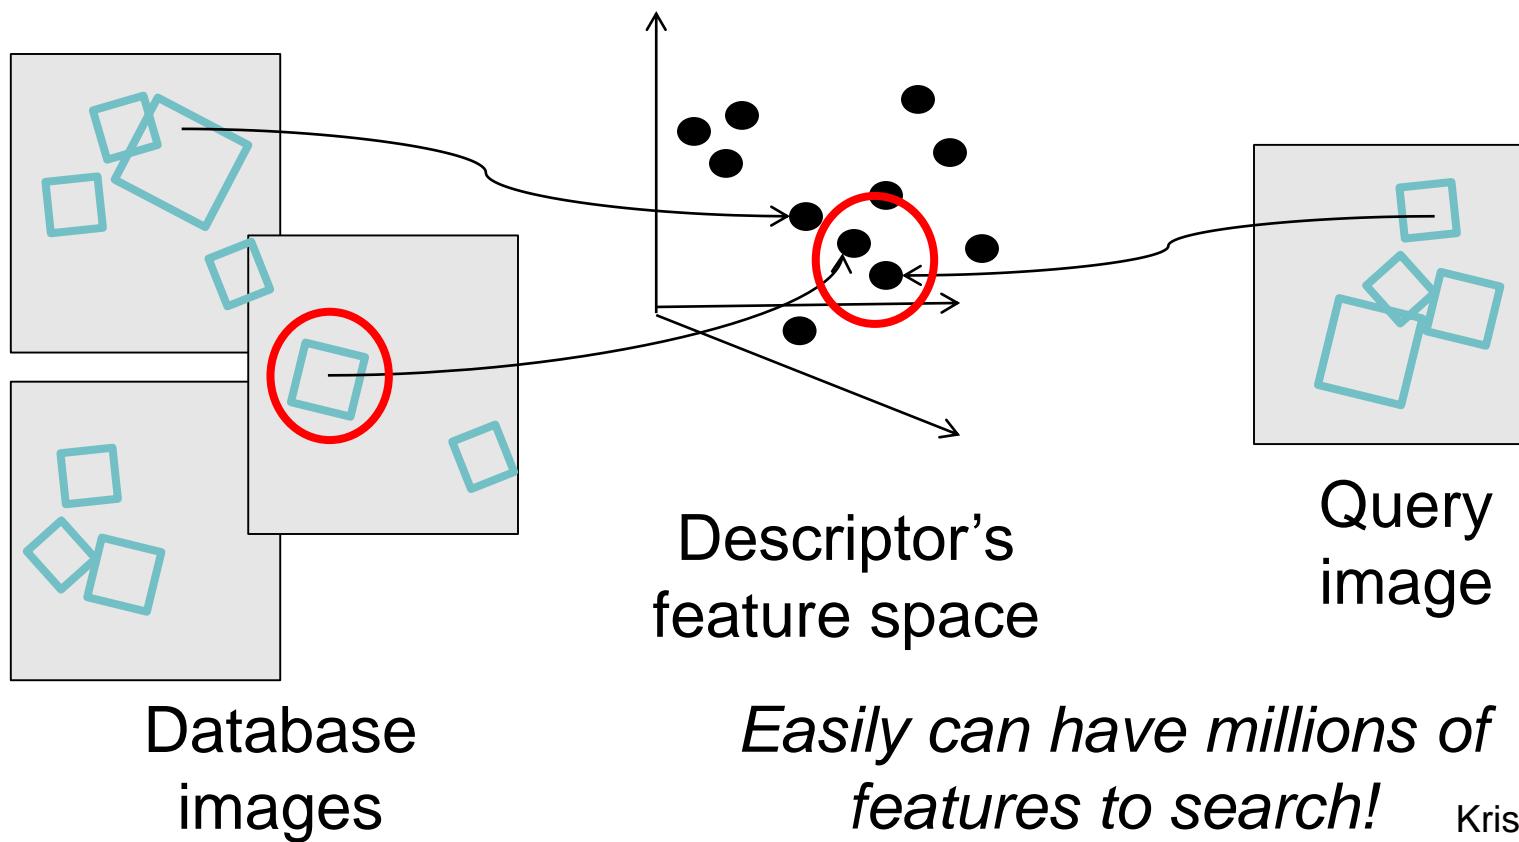
Indexing local features

- Each patch / region has a descriptor, which is a point in some high-dimensional feature space (e.g., SIFT)



Indexing local features

- When we see close points in feature space, we have similar descriptors, which indicates similar local content.



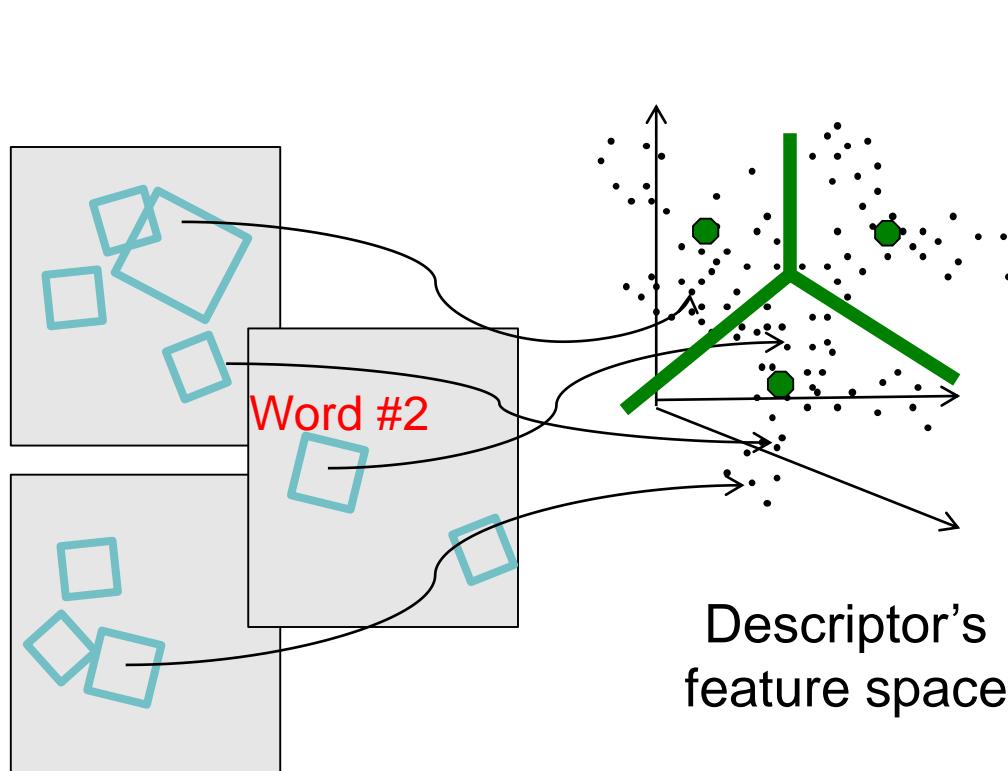
Indexing local features: inverted file index

Index	
"Along I-75," From Detroit to Florida; <i>inside back cover</i>	
"Drive I-95," From Boston to Florida; <i>inside back cover</i>	
1929 Spanish Trail Roadway; 101-102,104	
511 Traffic Information; 83	
A1A (Barrier Isl) - I-95 Access; 86	
AAA (and CAA); 83	
AAA National Office; 88	
Abbreviations, Colored 25 mile Maps; cover	
Exit Services; 196	
Travelogue; 85	
Africa; 177	
Agricultural Inspection Stns; 126	
Ah-Tah-Thi-Ki Museum; 160	
Air Conditioning, First; 112	
Alabama; 124	
Alachua; 132 County; 131	
Alafia River; 143	
Alapaha, Name; 126	
Alfred B Macay Gardens; 106	
Alligator Alley; 154-155	
Alligator Farm, St Augustine; 169	
Alligator Hole (definition); 157	
Alligator, Buddy; 155	
Alligators; 100,135,138,147,156	
Anastasia Island; 170	
Anhala; 108-109,146	
Apalachicola River; 112	
Appleton Mus of Art; 136	
Aquifer; 102	
Arabian Nights; 94	
Art Museum, Ringling; 147	
Aruba Beach Cafe; 183	
Aucilla River Project; 106	
Babcock-Web WMA; 151	
Bahia Mar Marina; 184	
Baker County; 99	
Barefoot Mallmen; 182	
Barge Canal; 137	
Bee Line Expy; 80	
Belz Outlet Mall; 89	
Bernard Castro; 136	
Big "I"; 165	
Big Cypress; 155,158	
Big Foot Monster; 105	
Billie Swamp Safari; 160	
Blackwater River SP; 117	
Blue Angels	
Butterfly Center, McGuire; 134	
CAA (see AAA)	
CCC, The; 111,113,115,135,142	
Ca' d'Zan; 147	
Caloosahatchee River; 152	
Name; 150	
Canaveral Natnl Seashore; 173	
Cannon Creek Airport; 130	
Canopy Road; 106,160	
Cape Canaveral; 174	
Castillo San Marcos; 169	
Cave Diving; 131	
Cayo Costa, Name; 150	
Celebration; 93	
Charlotte County; 149	
Charlotte Harbor; 150	
Chautauqua; 116	
Chipley; 114	
Name; 115	
Choctawatchee, Name; 115	
Circus Museum, Ringling; 147	
Citrus; 88,97,130,136,140,180	
CityPlace, W Palm Beach; 180	
City Maps,	
Ft Lauderdale Expwys; 194-195	
Jacksonville; 163	
Kissimmee Expwys; 192-193	
Miami Expressways; 194-195	
Orlando Expressways; 192-193	
Pensacola; 26	
Tallahassee; 191	
Tampa-St. Petersburg; 63	
St. Augustine; 191	
Civil War; 100,108,127,138,141	
Clearwater Marine Aquarium; 187	
Collier County; 154	
Collier, Barron; 152	
Colonial Spanish Quarters; 168	
Columbia County; 101,128	
Coquina Building Material; 165	
Corkscrew Swamp, Name; 154	
Cowboys; 95	
Crab Trap II; 144	
Cracker, Florida; 88,95,132	
Crosstown Expy; 11,35,98,143	
Cuban Bread; 184	
Dade Battlefield; 140	
Dade, Maj. Francis; 139-140,161	
Dania Beach Hurricane; 184	
Daniel Boone, Florida Walk; 117	
Daytona Beach; 172-173	
De Land; 87	
Driving Lanes; 85	
Duval County; 163	
Eau Gallie; 175	
Edison, Thomas; 152	
Eglin AFB; 116-118	
Eight Reale; 176	
Ellenton; 144-145	
Emanuel Point Wreck; 120	
Emergency Callboxes; 63	
Epiphytes; 142,148,157,159	
Escambia Bay; 119 Bridge (I-10); 119 County; 120	
Estero; 153	
Everglade,90,95,139-140,154-160	
Draining of; 156,181	
Wildlife MA; 160	
Wonder Gardens; 154	
Falling Waters SP; 115	
Fantasy of Flight; 95	
Fayer Dykes SP; 171	
Fires, Forest; 166	
Fires, Prescribed; 148	
Fisherman's Village; 151	
Flagler County; 171	
Flagler, Henry; 97,165,167,171	
Florida Aquarium; 186	
Florida,	
12,000 years ago; 187	
Cavern SP; 114	
Map of all Expressways; 2-3	
Mus of Natural History; 134	
National Cemetery ; 141	
Part of Africa; 177	
Platform; 187	
Sheriff's Boys Camp; 126	
Sports Hall of Fame; 130	
Sun 'n Fun Museum; 97	
Supreme Court; 107	
Florida's Turnpike (FTP); 178,189	
25 mile Strip Maps; 66	
Administration; 189	
Coin System; 190	
Exit Services; 189	
HEFT; 76,161,190	
History; 189	
Names; 189	
Service Plazas; 190	
Spur SR91; 76	
Ticket System; 190	
Toll Plazas; 190	
Ford, Henry; 152	

- For text documents, an efficient way to find all *pages* on which a *word* occurs is to use an index...
- We want to find all *images* in which a *feature* occurs.
- To use this idea, we'll need to map our features to “visual words”.

Visual words

- Map high-dimensional descriptors to tokens/words by quantizing the feature space



- Quantize via clustering, let cluster centers be the prototype “words”
- Determine which word to assign to each new image region by finding the closest cluster center.

Visual words

- Example: each group of patches belongs to the same visual word

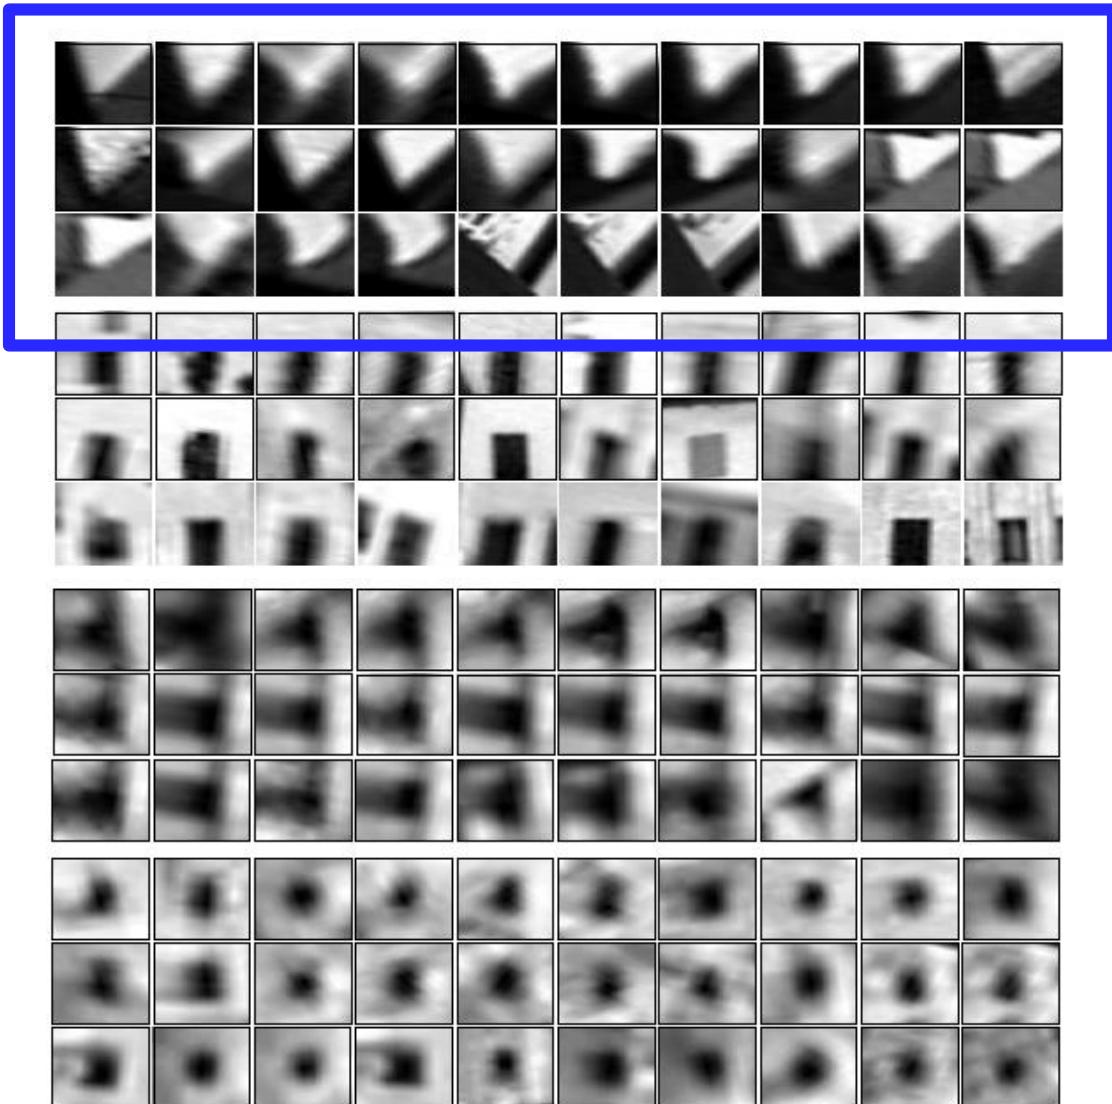
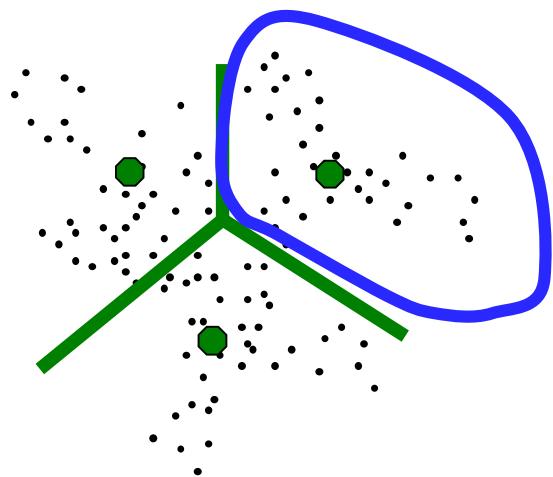


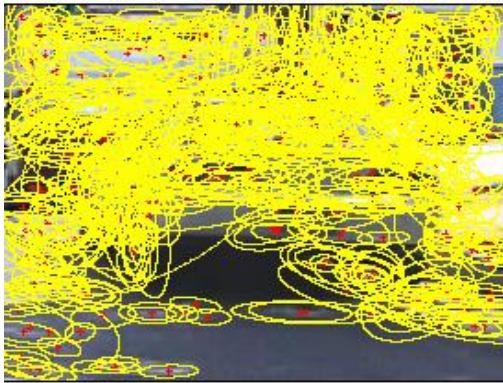
Figure from Sivic & Zisserman, ICCV 2003 Kristen Grauman

Visual vocabulary formation

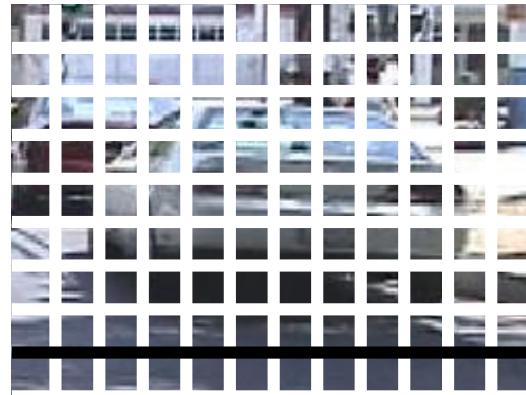
Issues:

- Vocabulary size, number of words
- Sampling strategy: where to extract features?
- Clustering / quantization algorithm
- Unsupervised vs. supervised
- What corpus provides features (universal vocabulary?)

Sampling strategies



Sparse, at interest points



Dense, uniformly



Randomly

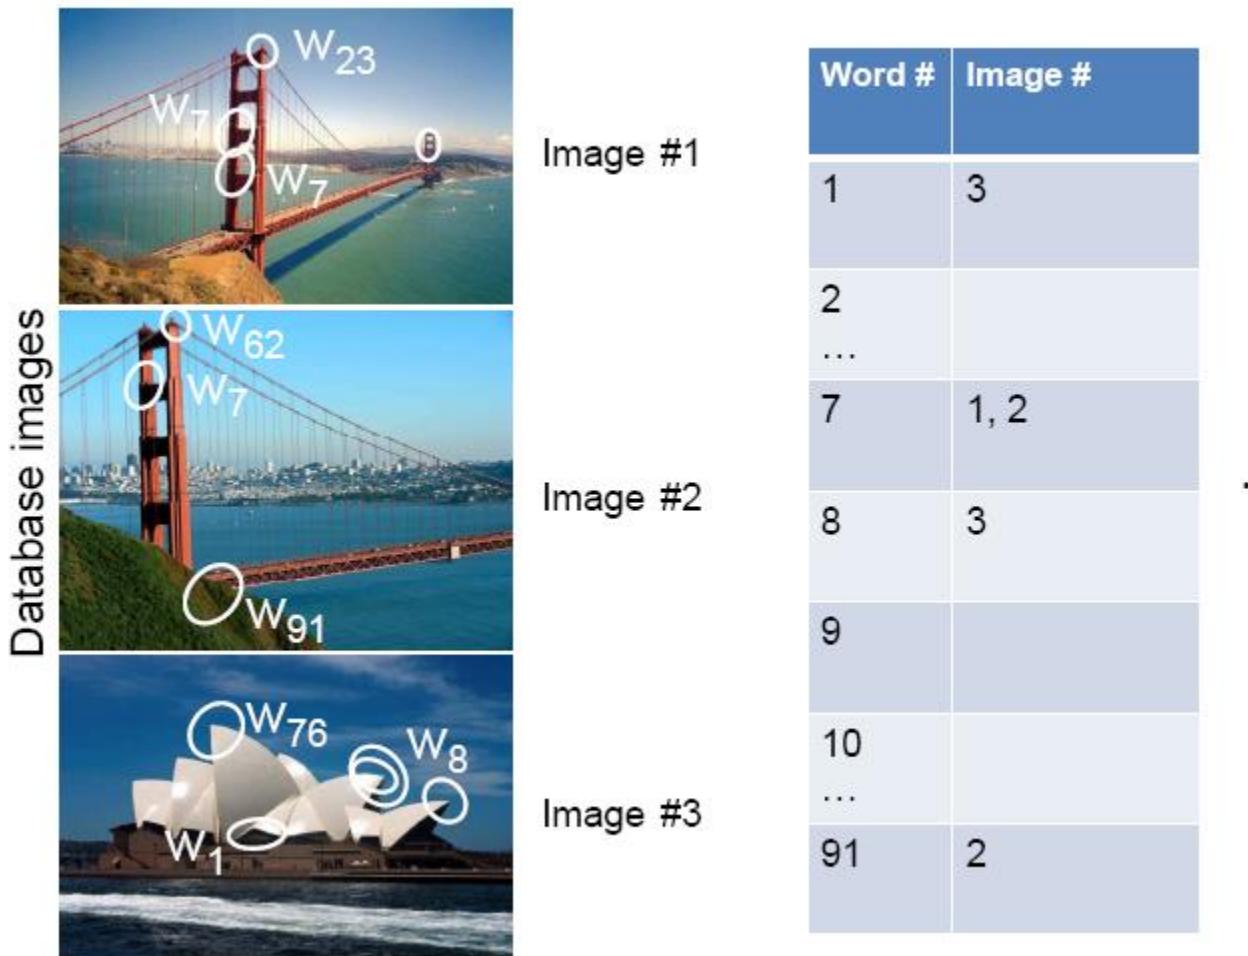


Multiple interest operators

- To find specific, textured objects, sparse sampling from interest points often more reliable.
- Multiple complementary interest operators offer more image coverage.
- For object categorization, dense sampling offers better coverage.

[See Nowak, Jurie & Triggs, ECCV 2006]

Inverted file index



- Database images are loaded into the index mapping words to image numbers

Inverted file index



New query image

Word #	Image #
1	3
2	
7	1, 2
8	3
9	
10	
...	
91	2



- New query image is mapped to indices of database images that share a word.

Inverted file index

- Key requirement for inverted file index to be efficient: sparsity
- If most pages/images contain most words then you're no better off than exhaustive search.
 - Exhaustive search would mean comparing the word distribution of a query versus every page.

Instance recognition: remaining issues

- How to summarize the content of an entire image? And gauge overall similarity?
- How large should the vocabulary be? How to perform quantization efficiently?
- Is having the same set of visual words enough to identify the object/scene? How to verify spatial agreement?
- How to score the retrieval results?